

Algorithmic Human-Robot Interaction

Task Planning III and Project Pitches

CSCI 7000

Prof. Brad Hayes

Computer Science Department

University of Colorado Boulder

Papers for Thursday:

Need 1 PRO (10m) and 1 CON (5m) speaker each

Designing Robot Learners that Ask Good Questions

Maya Cakmak and Andrea Thomaz

PRO presenter: Lakhan Kamireddy

CON presenter: Dhanendra Soni

Anticipating human actions for collaboration in the presence of task and sensor uncertainty

Kelsey Hawkins et al.

PRO presenter: Ian Loegfren

CON presenter: Nishank Sharma


Last Time...

Keyframe and Trajectory Demonstrations




Learning Kitchen Tasks
from
Hybrid Demonstrations


Trajectories and Keyframes for Kinesthetic Teaching: A Human-Robot Interaction Perspective



Trajectory
Demonstration




Keyframe
Demonstration




Hybrid
Demonstration

Sample demonstrations of the letter P in 2D

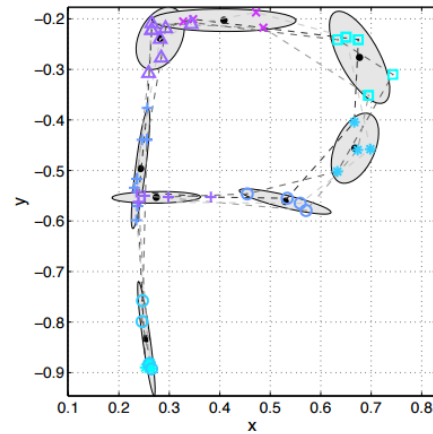
Trajectory Conversion




Continuous
trajectories in 2D



Data converted
to keyframes



Clustering of keyframes
and the sequential
pose distributions



Learned model
trajectory

Trajectory Conversion: Forward-Inverse Relaxation Model





Figure 3: An algorithm for extracting via-points.

- Fifth order splines used between positions to minimize jerk, using position, velocity, and acceleration per keyframe to compute the spline unknowns.
- Keyframes assume zero velocity/acceleration per point
- Trajectory demonstrations use the means from cluster centers.

Aligning Multiple Demonstrations


Keyframe demonstrations




Temporal alignment

a b c d
x y y z


Clustering



Planning human-aware motions using a sampling-based costmap planner



Extending this work with human motion models



STRIPS: Language for Classical Planning

- A **problem** in Strips is a tuple $P = \langle F, O, I, G \rangle$:
 - ▷ F stands for set of all **atoms** (boolean vars)
 - ▷ O stands for set of all **operators** (actions)
 - ▷ $I \subseteq F$ stands for **initial situation**
 - ▷ $G \subseteq F$ stands for **goal situation**
- Operators $o \in O$ **represented** by
 - ▷ the **Add** list $Add(o) \subseteq F$
 - ▷ the **Delete** list $Del(o) \subseteq F$
 - ▷ the **Precondition** list $Pre(o) \subseteq F$
- Pickup(X)
 - **P**: $\text{grip}(\emptyset) \wedge \text{clear}(X) \wedge \text{ontable}(X)$
 - **A**: $\text{grip}(X)$
 - **D**: $\text{onTable}(X) \wedge \text{grip}(\emptyset)$

Planning with Markov Decision Processes

MDPs are **fully observable, probabilistic** state models:

- a state space S
 - initial state $s_0 \in S$
 - a set $G \subseteq S$ of goal states
 - actions $A(s) \subseteq A$ applicable in each state $s \in S$
 - **transition probabilities** $P_a(s'|s)$ for $s \in S$ and $a \in A(s)$
 - action costs $c(a, s) > 0$
-
- **Solutions** are **functions (policies)** mapping states into actions
 - **Optimal** solutions minimize **expected cost** to goal

Partially Observable MDPs (POMDPs)

- Traditional MDPs are defined with:

- States — $S = \{(0,0), (0,1), \dots\}$
- Actions — $A = \{\text{move_north}, \dots\}$
- Rewards — $R(s,a,s') \rightarrow \text{Reward}$
- Transition Probabilities — $T(s,a,s') \rightarrow P(s' \mid s, a)$

- Now we have to add:

- Observation set — $O = \{o_1, o_2, \dots\}$
- Observation prob. — $\Omega = P(o_1, \dots, o_i \mid S)$


- Also have to augment:

- Current State (*now belief*) — $B = [0.1, 0.6, 0.2, 0.1, \dots]$
- Policy (*no longer $S \rightarrow A$*) — $\pi: B \rightarrow A$

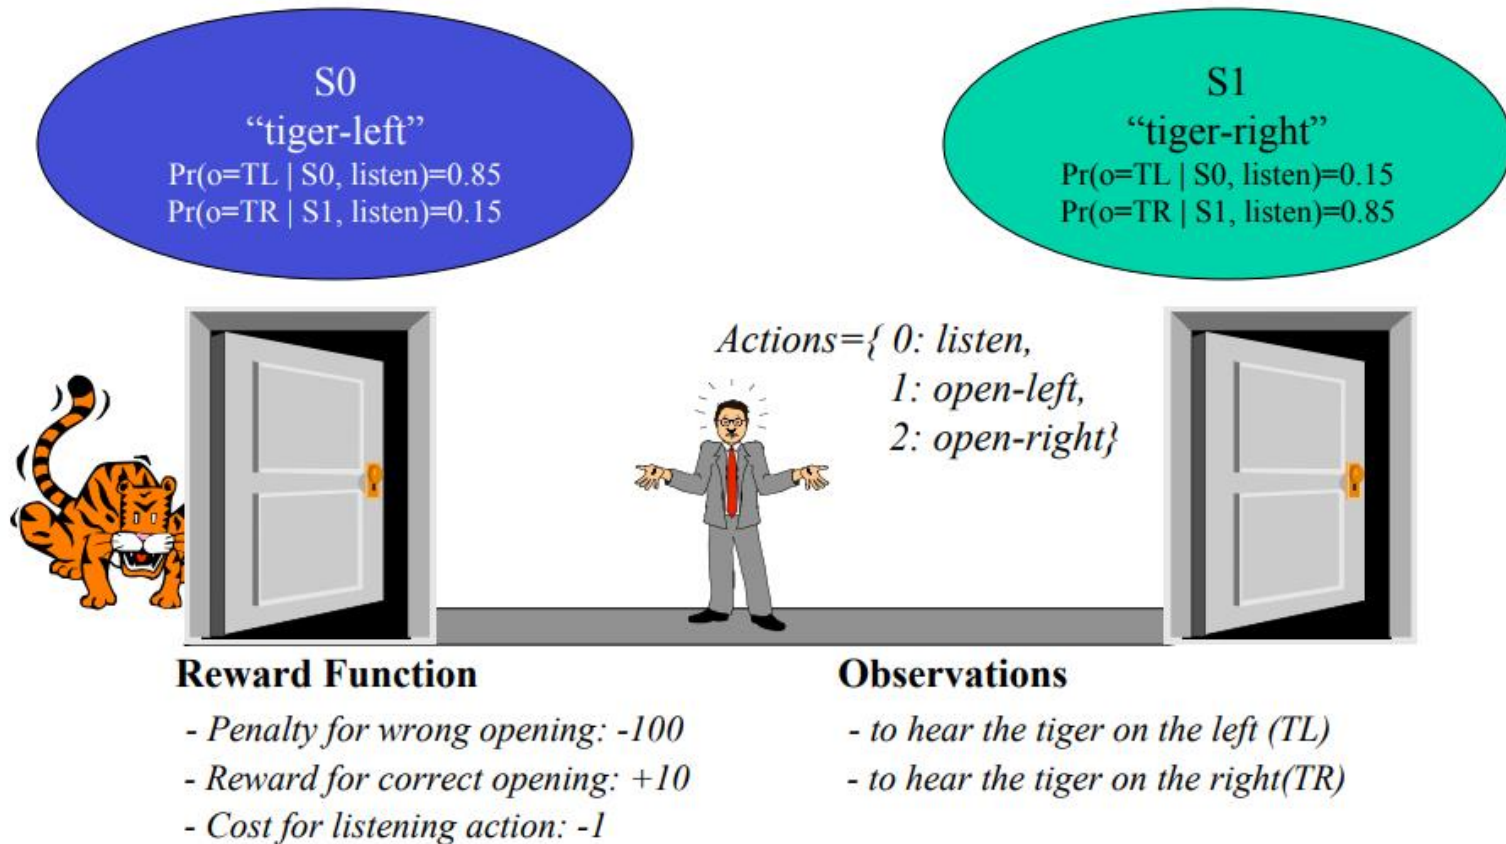
POMDPs in Action



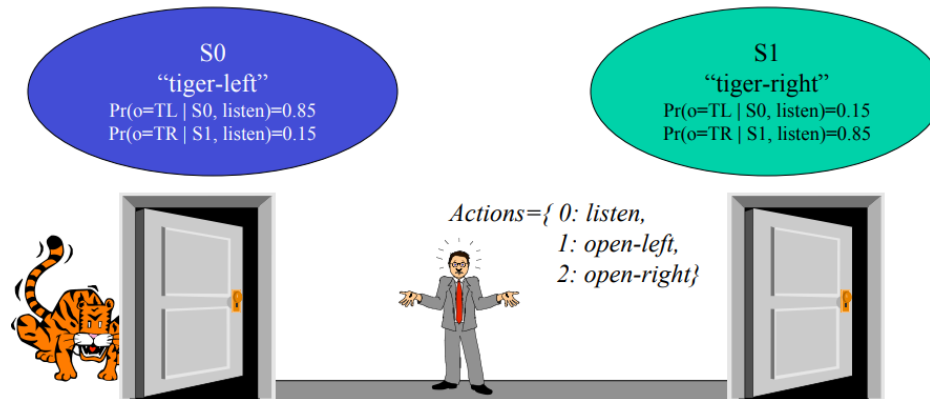
POMDPs in Action



POMDP: The Tiger Problem



POMDP: The Tiger Problem



Prob. (LISTEN)	Tiger: left	Tiger: right
Tiger: left	1.0	0.0
Tiger: right	0.0	1.0

Prob. (LEFT)	Tiger: left	Tiger: right
Tiger: left	0.5	0.5
Tiger: right	0.5	0.5

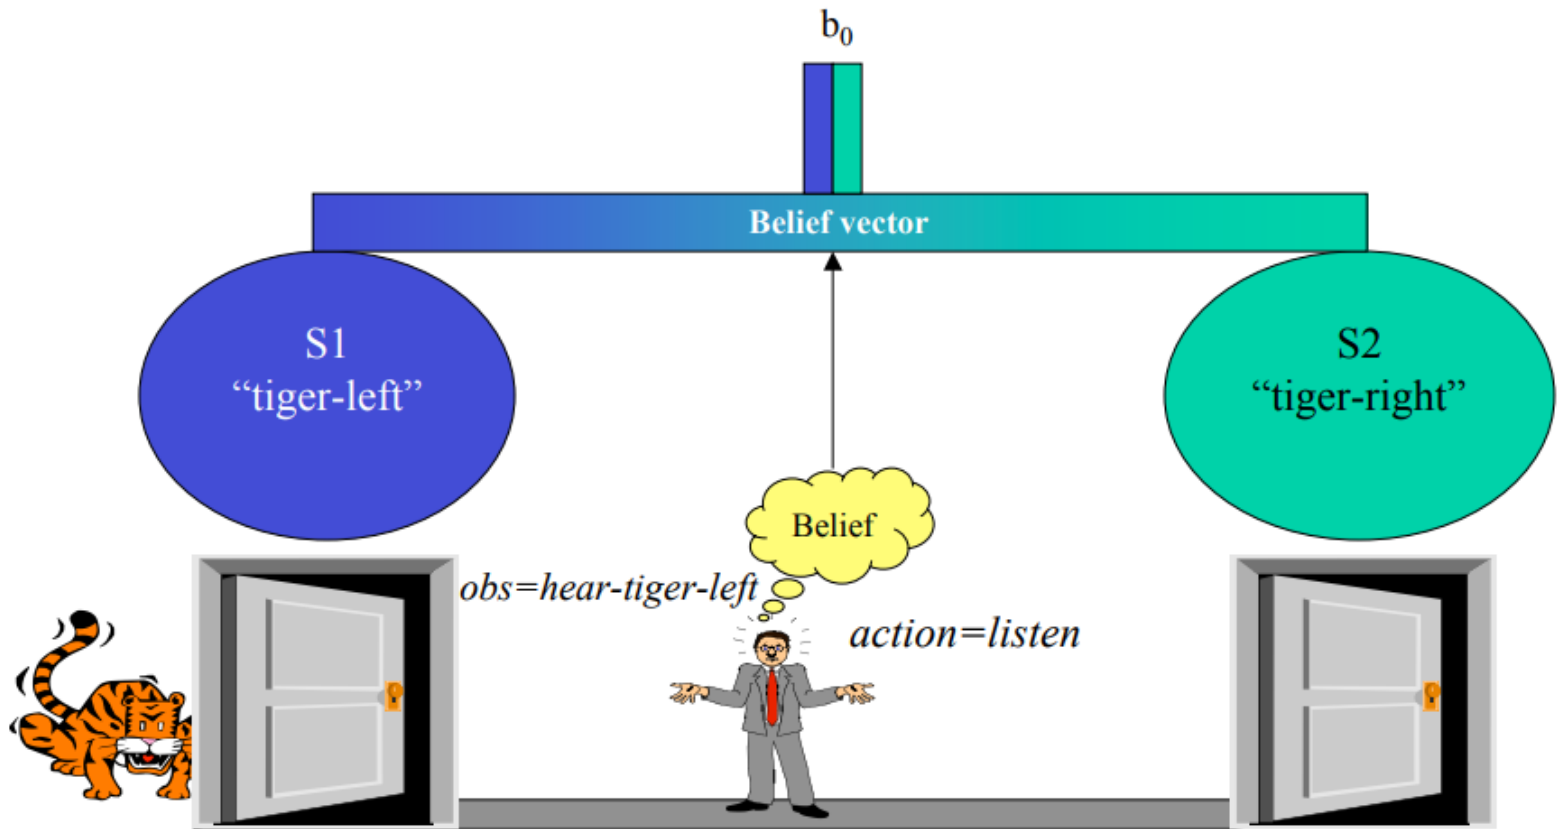
Prob. (RIGHT)	Tiger: left	Tiger: right
Tiger: left	0.5	0.5
Tiger: right	0.5	0.5

Prob. (LISTEN)	O: TL	O: TR
Tiger: left	0.85	0.15
Tiger: right	0.15	0.85

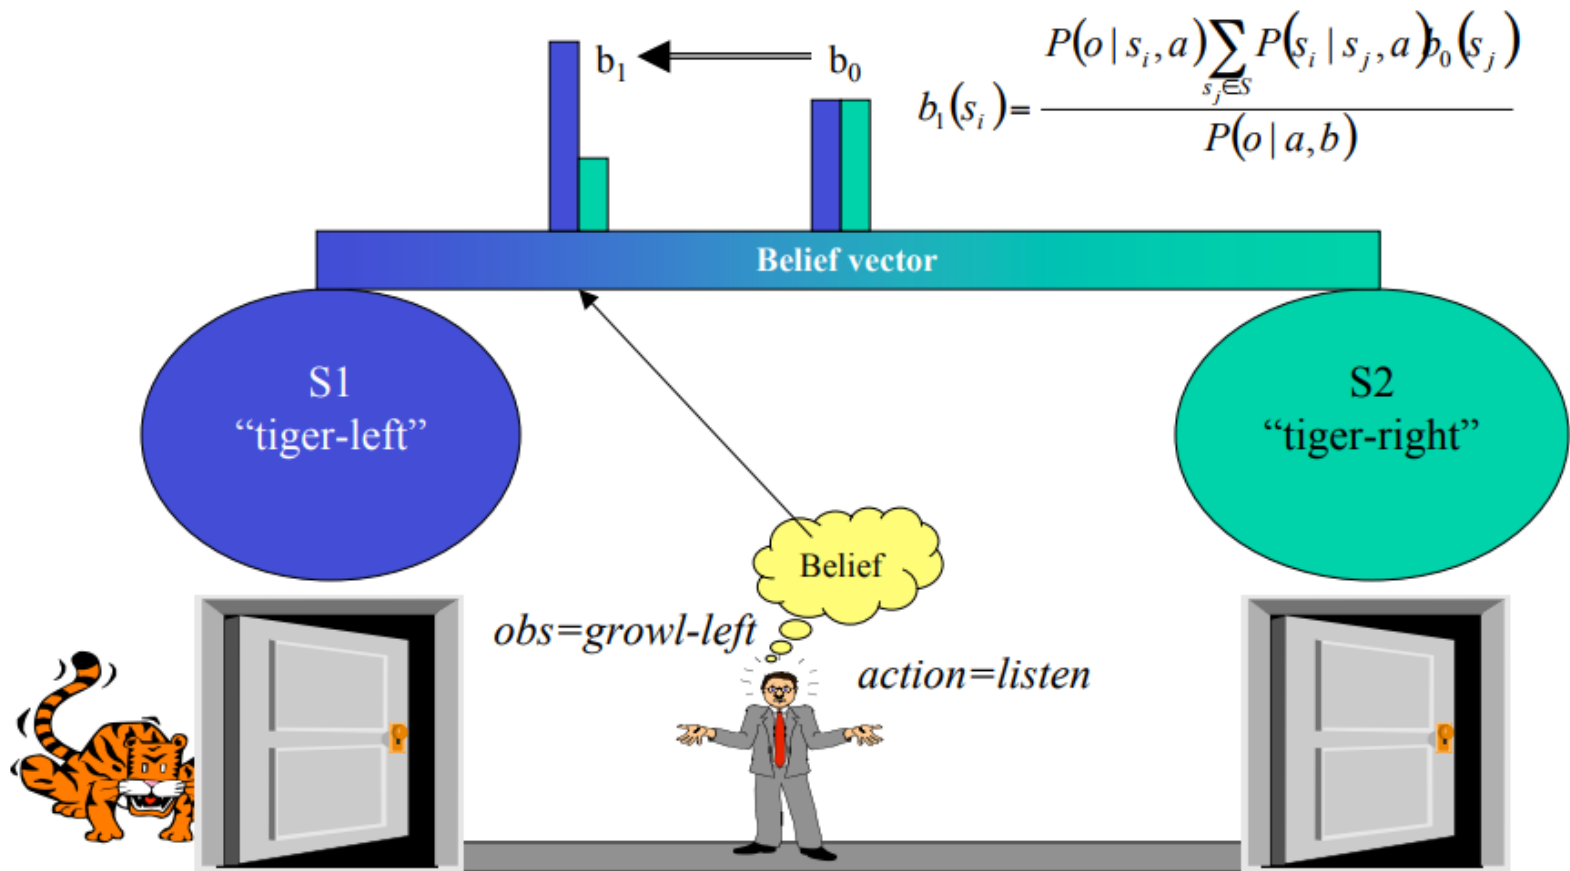
Prob. (LEFT)	O: TL	O: TR
Tiger: left	0.5	0.5
Tiger: right	0.5	0.5

Prob. (LEFT)	O: TL	O: TR
Tiger: left	0.5	0.5
Tiger: right	0.5	0.5

POMDP: The Tiger Problem




POMDP: The Tiger Problem




Still not feasible to exhaustively check all of our hypotheses!

- But we can still make progress!
- Need a mechanism to enable us to mostly pursue *reasonably likely* hypotheses.
- Example domain: Localization!



Localization Domain

- **Latent variable:** Horizontal aircraft position
- **Observable variable:** Vertical aircraft position (with some error)
- **Known information:** Airspeed (with some error), Map
- **Goal:** Use repeated observations to find true horizontal position




Particle Filter

- **Latent variable:** Horizontal aircraft position
- **Observable variable:** Vertical aircraft position (altitude)
- **Known information:** Airspeed, Map
- **Goal:** Use repeated observations to find true horizontal position

Main Idea:

Generate lots of hypotheses
and let the observations
determine their likelihood.




Particle Filter: Initial Hypotheses

- **Latent variable:** Horizontal aircraft position
- **Observable variable:** Vertical aircraft position (altitude)
- **Known information:** Airspeed, Map
- **Goal:** Use repeated observations to find true horizontal position

Main Idea:

Generate lots of hypotheses and let the observations determine their likelihood.




Particle Filter: Likelihood Evaluation

- **Latent variable:** Horizontal aircraft position
- **Observable variable:** Vertical aircraft position (altitude)
- **Known information:** Airspeed, Map
- **Goal:** Use repeated observations to find true horizontal position

Main Idea:

Generate lots of hypotheses and let the observations determine their likelihood.




Particle Filter: Weighting Particles

- **Latent variable:** Horizontal aircraft position
- **Observable variable:** Vertical aircraft position (altitude)
- **Known information:** Airspeed, Map
- **Goal:** Use repeated observations to find true horizontal position

Main Idea:

Generate lots of hypotheses and let the observations determine their likelihood.




Particle Filter: Exhausting Particles

- **Latent variable:** Horizontal aircraft position
- **Observable variable:** Vertical aircraft position (altitude)
- **Known information:** Airspeed, Map
- **Goal:** Use repeated observations to find true horizontal position

Main Idea:

Generate lots of hypotheses and let the observations determine their likelihood.

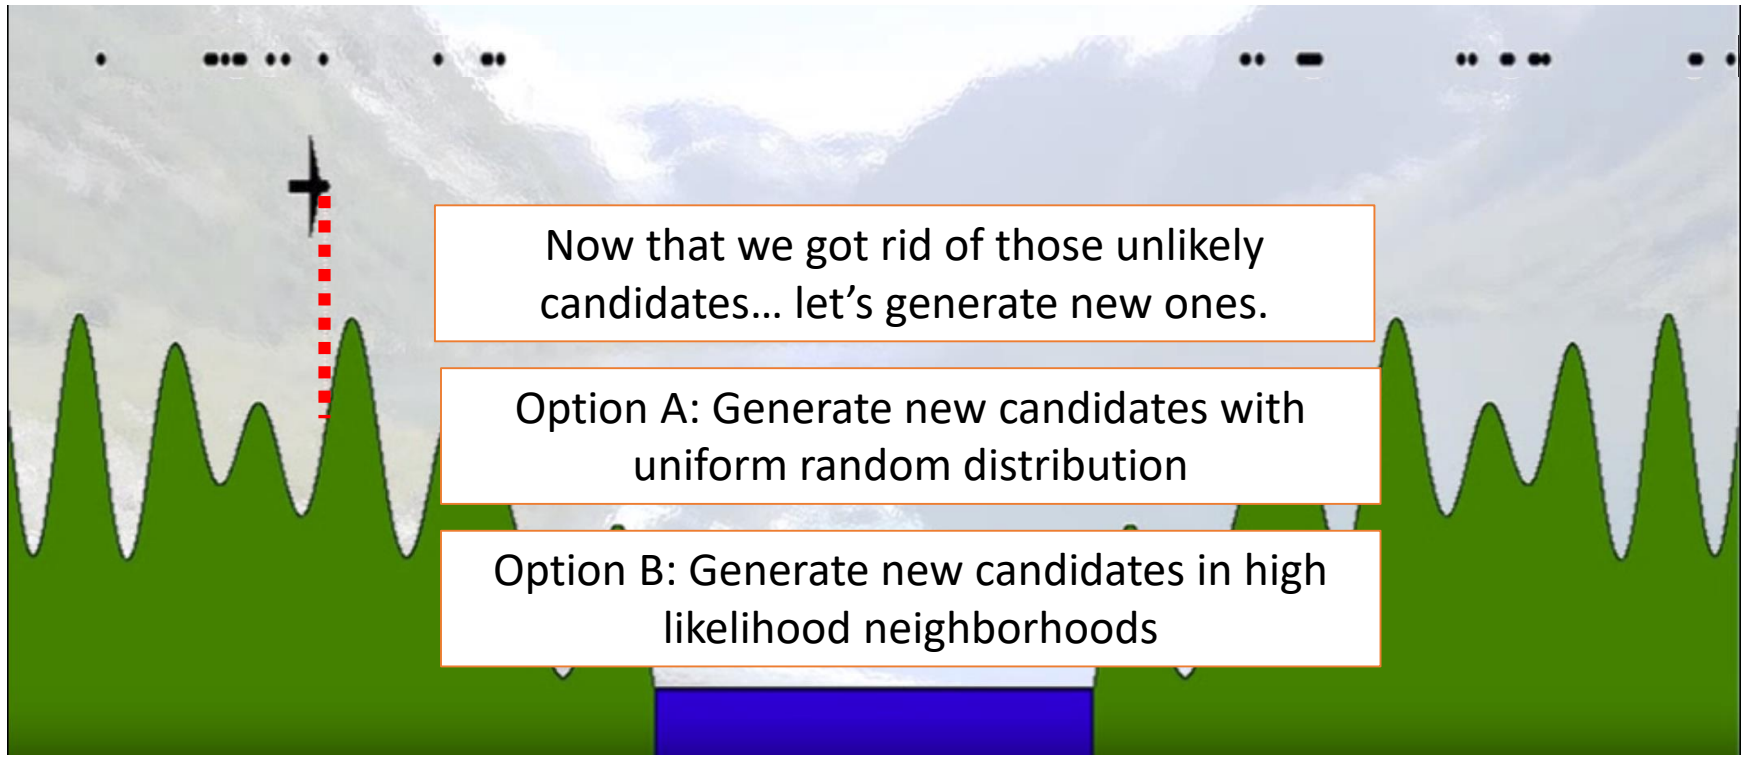


Particle Filter: Resampling Particles

- **Latent variable:** Horizontal aircraft position
- **Observable variable:** Vertical aircraft position (altitude)
- **Known information:** Airspeed, Map
- **Goal:** Use repeated observations to find true horizontal position

Main Idea:

Generate lots of hypotheses and let the observations determine their likelihood.



Now that we got rid of those unlikely candidates... let's generate new ones.

Option A: Generate new candidates with uniform random distribution


Option B: Generate new candidates in high likelihood neighborhoods

Particle Filter: Resampling Particles

- **Latent variable:** Horizontal aircraft position
- **Observable variable:** Vertical aircraft position (altitude)
- **Known information:** Airspeed, Map
- **Goal:** Use repeated observations to find true horizontal position

Main Idea:

Generate lots of hypotheses and let the observations determine their likelihood.




Particle Filter: Advance 1 Timestep

- **Latent variable:** Horizontal aircraft position
- **Observable variable:** Vertical aircraft position (altitude)
- **Known information:** Airspeed, Map
- **Goal:** Use repeated observations to find true horizontal position

Main Idea:

Generate lots of hypotheses
and let the observations
determine their likelihood.




Particle Filter: Update Weights

- **Latent variable:** Horizontal aircraft position
- **Observable variable:** Vertical aircraft position (altitude)
- **Known information:** Airspeed, Map
- **Goal:** Use repeated observations to find true horizontal position

Main Idea:

Generate lots of hypotheses and let the observations determine their likelihood.




Particle Filter: Exhaust/Resample

- **Latent variable:** Horizontal aircraft position
- **Observable variable:** Vertical aircraft position (altitude)
- **Known information:** Airspeed, Map
- **Goal:** Use repeated observations to find true horizontal position

Main Idea:

Generate lots of hypotheses and let the observations determine their likelihood.




Particle Filter: Exhaust/Resample

- **Latent variable:** Horizontal aircraft position
- **Observable variable:** Vertical aircraft position (altitude)
- **Known information:** Airspeed, Map
- **Goal:** Use repeated observations to find true horizontal position

Main Idea:

Generate lots of hypotheses
and let the observations
determine their likelihood.




Particle Filter: Advance 1 Timestep

- **Latent variable:** Horizontal aircraft position
- **Observable variable:** Vertical aircraft position (altitude)
- **Known information:** Airspeed, Map
- **Goal:** Use repeated observations to find true horizontal position

Main Idea:

Generate lots of hypotheses
and let the observations
determine their likelihood.




Particle Filter: Future Iterations...

- **Latent variable:** Horizontal aircraft position
- **Observable variable:** Vertical aircraft position (altitude)
- **Known information:** Airspeed, Map
- **Goal:** Use repeated observations to find true horizontal position

Main Idea:

Generate lots of hypotheses
and let the observations
determine their likelihood.




More Particle Filtering in Action!

Real-Time Particle Filter Localization Demo

Stata Basement Loop
Instructor Solution by Corey Walsh
6.141 Spring 2017
Lab 5

Markov Chains


- Finite number of discrete states
- Probabilistic transitions between states
- Next state determined only by the current state
 - This is the Markov property



Rewards: $S1 = 10$, $S2 = 0$

Hidden Markov Model

- Finite number of discrete states
- Probabilistic transitions between states
- Next state determined only by the current state
- We're unsure which state we're in
 - The current states emits an observation



Rewards: $S1 = 10$, $S2 = 0$


Do not know state:

S1 emits O1 with prob 0.75

S2 emits O2 with prob 0.75

Markov Decision Process


- Finite number of discrete states
- Probabilistic transitions between states **and** controllable actions in each state
- Next state determined only by the current state **and** current action
 - This is still the Markov property



Rewards: $S1 = 10$, $S2 = 0$

Partially Observable Markov Decision Process

- Finite number of discrete states
- Probabilistic transitions between states and controllable actions
- Next state determined only by the current state and current action
- We're unsure which state we're in
 - The current state emits observations



Rewards: $S1 = 10$, $S2 = 0$

Do not know state:

S1 emits O1 with prob 0.75

S2 emits O2 with prob 0.75

Markov Model Chart

**Do we have control over the state transitions?
(Are we picking which actions are executed)**

**Are the states
completely
observable?**

	NO	YES
YES	Markov Chain	MDP
NO	HMM	POMDP



Final Project Pitches

In two minutes or less:

What new capability are you introducing / phenomenon are you exploring?

What's hard about this now?

How will you evaluate whether your approach works or not?

