

The background is a collage of images related to human-robot interaction. It includes a person wearing a sensor vest, a robotic arm with a yellow glove, a person working with a robotic arm, and a person sitting at a desk with a laptop. The text is overlaid on this collage.

Algorithmic Human-Robot Interaction

Modeling with GMMs and HMMs for Activity Recognition

CSCI 7000

Prof. Brad Hayes

University of Colorado Boulder



HUMAN ROBOT COLLABORATION

Last Time...



Papers for Thursday 2/28: Natural Language Understanding

Robust Robot Learning from Demonstration and Skill Repair Using Conceptual Constraints – Mueller et al.

Pro: Jack Kawell

Con: Matthew Luebbers

Accurately and Efficiently Interpreting Human-Robot Instructions of Varying Granularities – Arumugam et al.

Pro: Karthik Palavalli

Con: Ian Loefgren

Introduction to Machine Learning

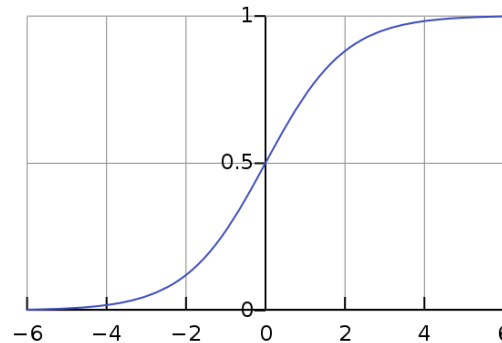
- Regression: How much is this house worth?
- Classification: Is this a photo of a dog or ice cream?



Linear Regression to Logistic Regression

- Linear Regression gives us a continuous-valued function approximation
 - Models relationship between scalar dependent variable y and one or more variables X
- Logistic regression allows us to approximate **categorical** data
 - Pick a model function that squashes values between 0 and 1

$$F(x) = \frac{1}{1 + e^{-x}}$$



- Apply it to a familiar function: $g(X) = \beta_0 + \beta_1 x + \epsilon$

$$P(Y = 1) = F(g(x)) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 * x)}}$$

Training and Testing Your Algorithms

- Partition your data into TRAIN, VALIDATE, and TEST
- Train your model on TRAIN
- Evaluate your model on VALIDATE
- TRAIN and VALIDATE can be swapped around
- You only get to run on TEST once, ever!



Types of Learning

Supervised Learning

- Given a dataset of (X, y) pairs
 - X : Vector representing a data point
 - y : Label or value that is the correct answer for $f(X) = y$
- “You guess, I tell you the correct answer, you update, repeat”

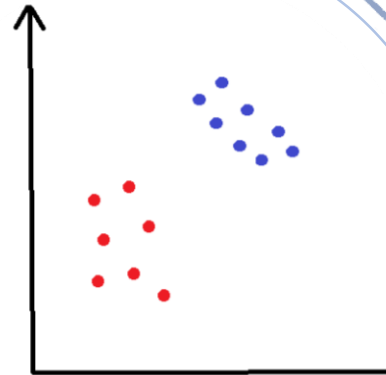
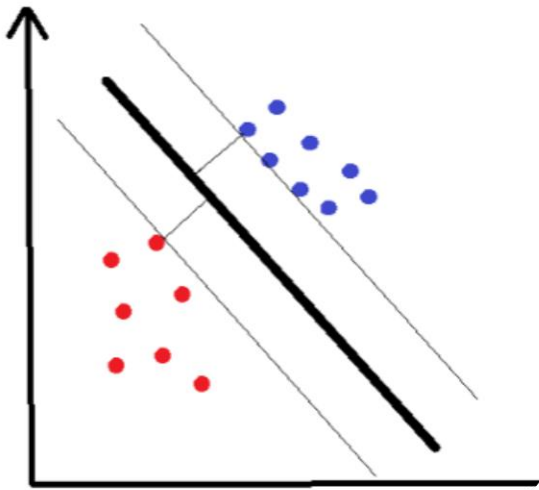
Reinforcement Learning

- Given some reward function R
 - $R(s, a, s')$ gives the value of moving from state s to s' via action a
- “You guess, I tell you if you’re on the right track (sometimes)”

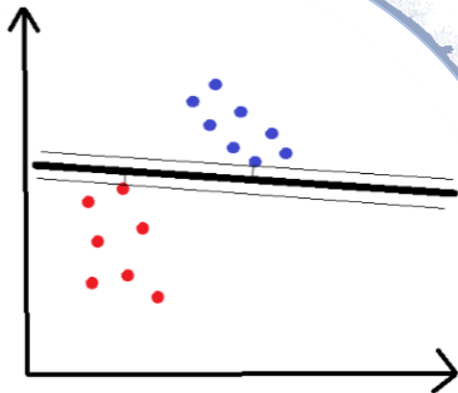
Unsupervised Learning

- Given a dataset of X ’s
 - X : Vector representing a data point
 - No labels (answers) given!
- “You guess, I have no feedback to give you. Good luck!”

Supervised Learning: Support Vector Machines



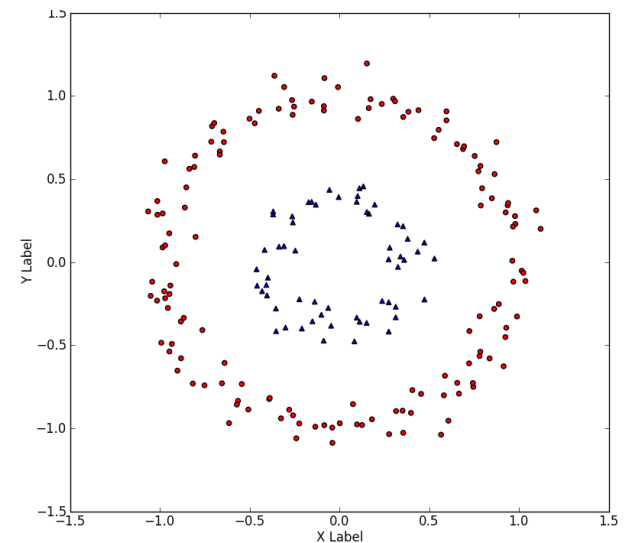
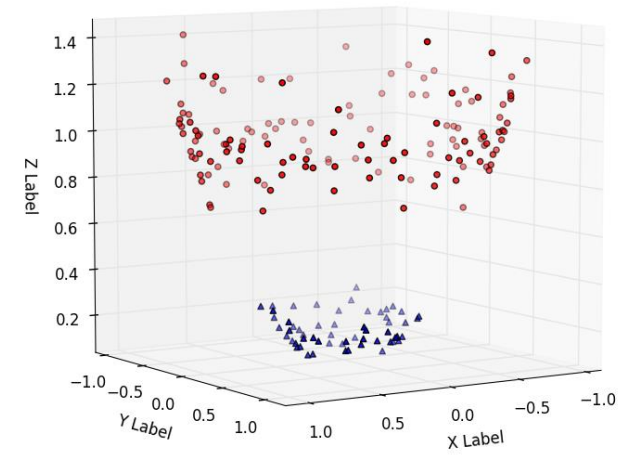
- Is this dot red or blue?
- Classifier will draw a separating line (or hyperplane)



Supervised Learning: Support Vector Machines

What happens if the data doesn't separate cleanly?

- Add a cost for misclassified examples and let the optimizer take care of it!
- Add more dimensions to the data!
 - $x^2, y^2, x^2 + y^2, \cos(x), xy$, etc.



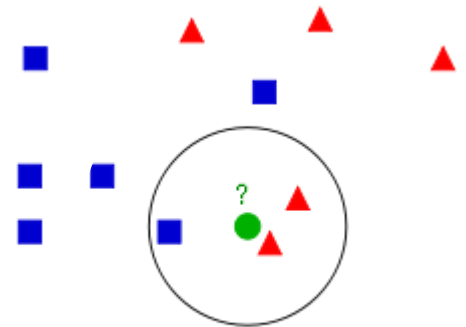
Supervised Learning: Support Vector Machines

Practical details:

- Scikit-Learn (sklearn) Python Package has excellent documentation
 - <http://scikit-learn.org/stable/modules/svm.html>
- For easier problems, can pretty much use it out of the box to get decent results

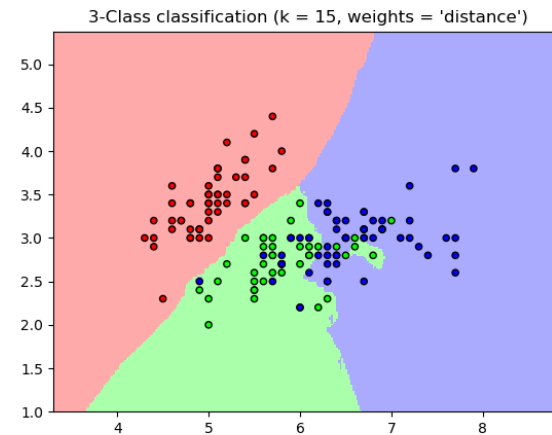
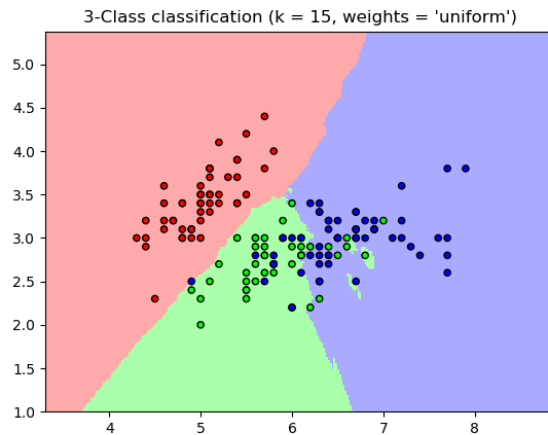
K-Nearest Neighbor

- Can be used for classification or regression
- Simple idea:
 - For a given data point p , find the K nearest labeled points
 - Assign the majority label to p
- Caveats:
 - The order that you label points can matter!
 - Slow! Lots of comparisons required.
 - Choosing 'K' has a big impact



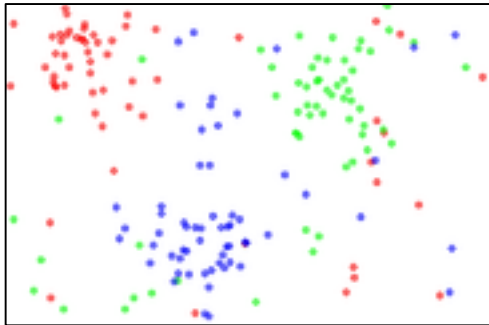
Weighted K-Nearest Neighbor

Weight each sample's influence by how far away it is from p

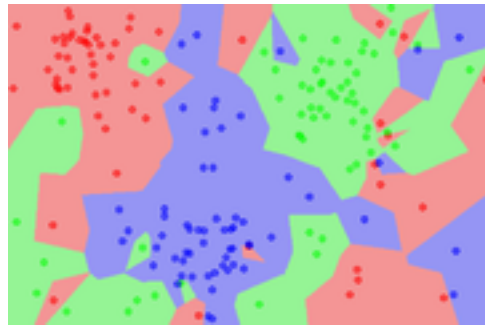


Reduced K-Nearest Neighbor

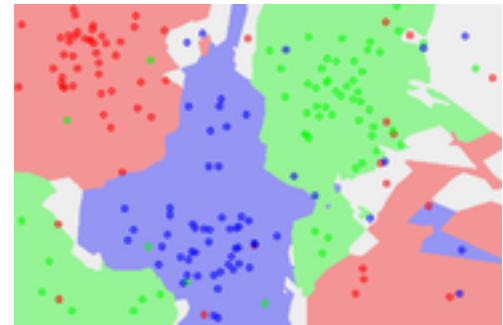
Dataset



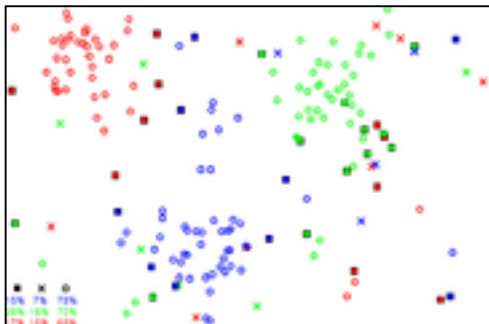
1-NN Classification



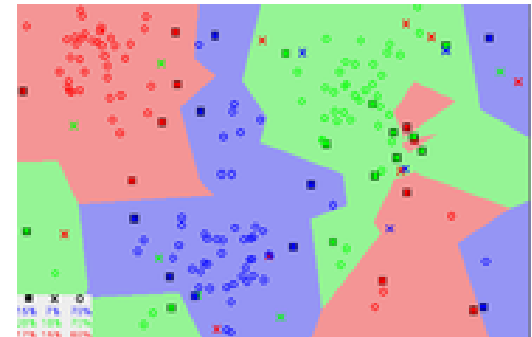
5-NN Classification



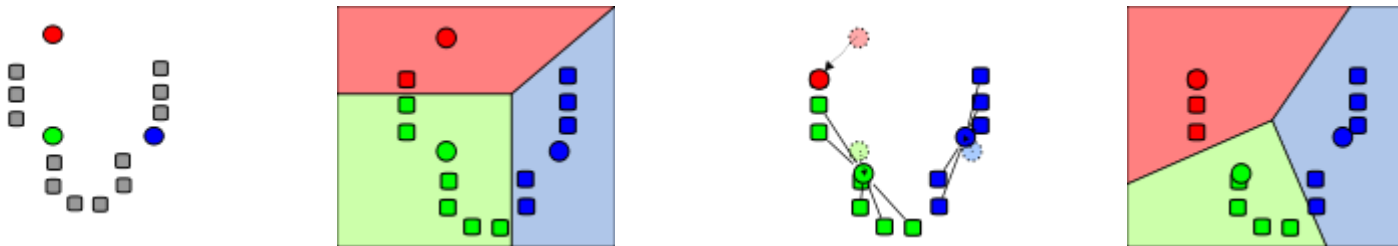
Reduced Dataset



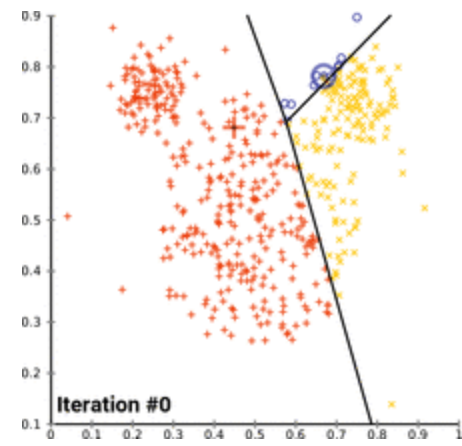
1-NN Classification



Unsupervised Learning: K-Means Clustering



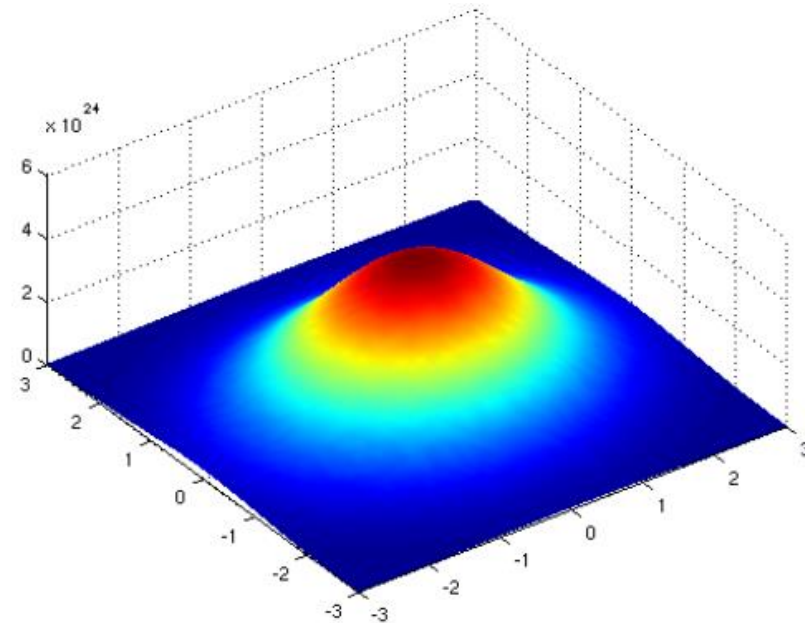
1. Randomly initialize k clusters at random positions.
2. Classify every data point as belonging to a cluster by Euclidean distance measurement (closest cluster wins)
3. Calculate centroid of each cluster. Relocate cluster to centroid position.
4. Repeat 2-3 until centroids converge.



Gaussian Distribution

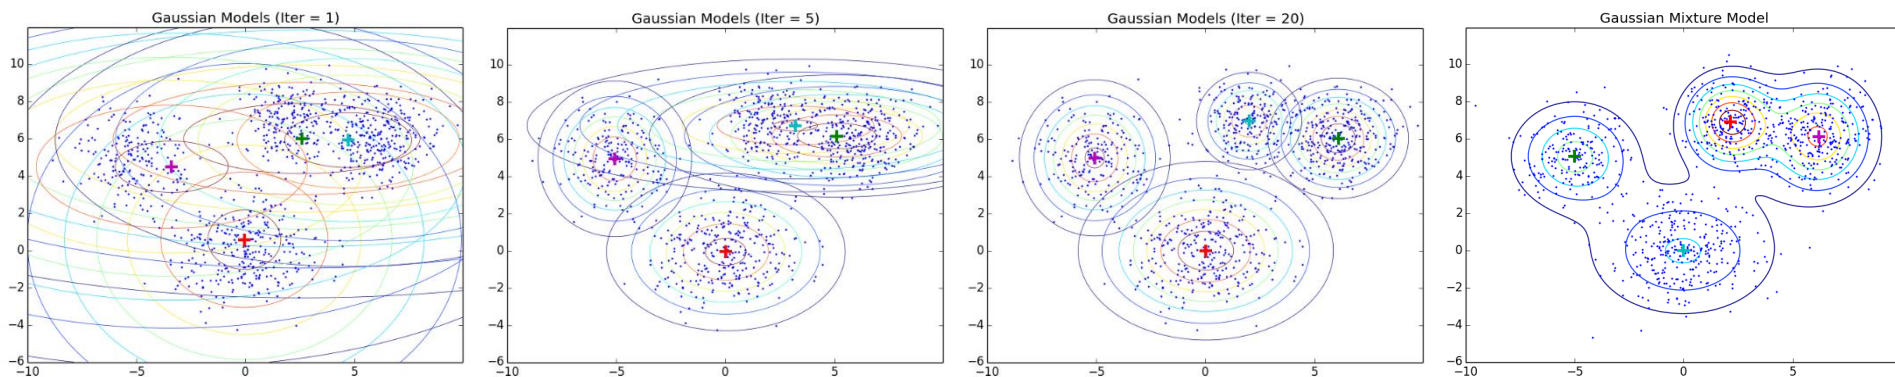
$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

$$\mathcal{N}(\mathbf{x}|\mu, \mathbf{\Sigma}) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\mathbf{\Sigma}|^{1/2}} e^{-\frac{1}{2}(\mathbf{x}-\mu)^T \mathbf{\Sigma}^{-1}(\mathbf{x}-\mu)}$$

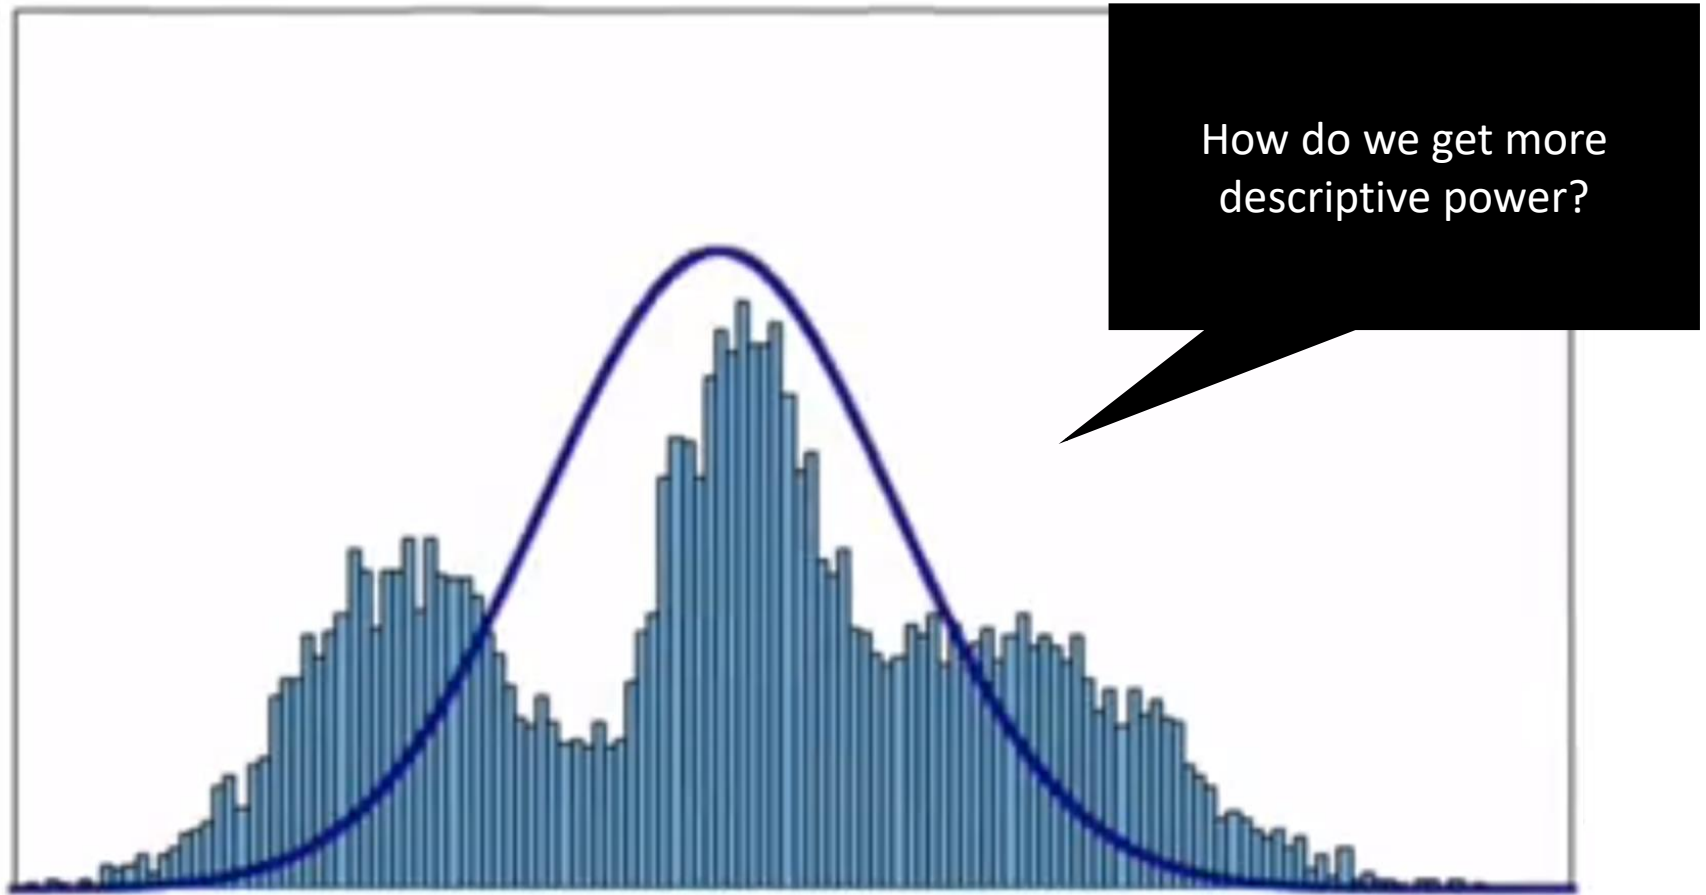


Gaussian Mixture Model (GMM)

- Model that represents a distribution that data are drawn from
 - Lets us classify existing data, and generate “plausible” new data!
- Generalization of k-means clustering to incorporate information about the covariance of the data

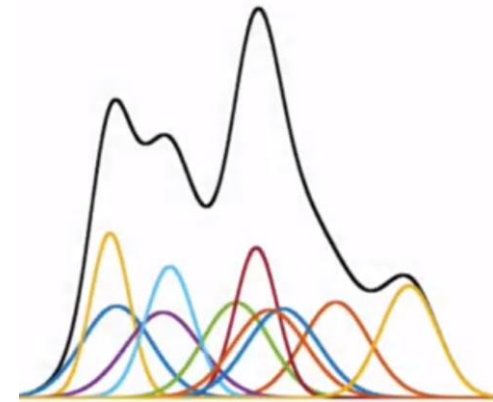


Modeling with one Gaussian



Mixture: Sum of Gaussians

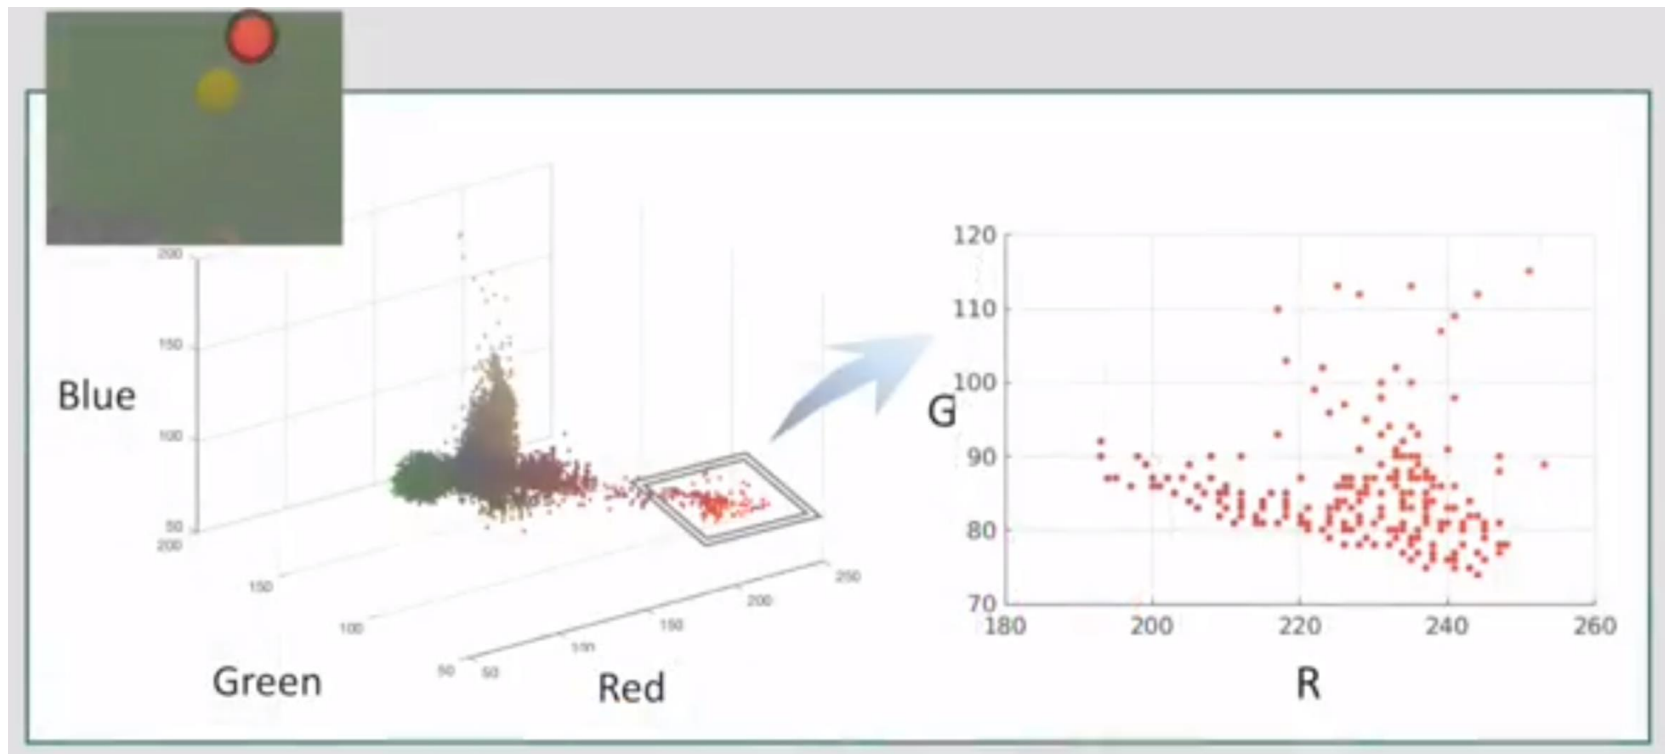
$$p(x) = \sum_{k=1}^K w_k * g_k(x | \mu_k, \Sigma_k)$$



g_k is a Gaussian distribution with mean μ_k and covariance matrix Σ_k
 w_k is the mixing coefficient of a particular Gaussian (its weight)

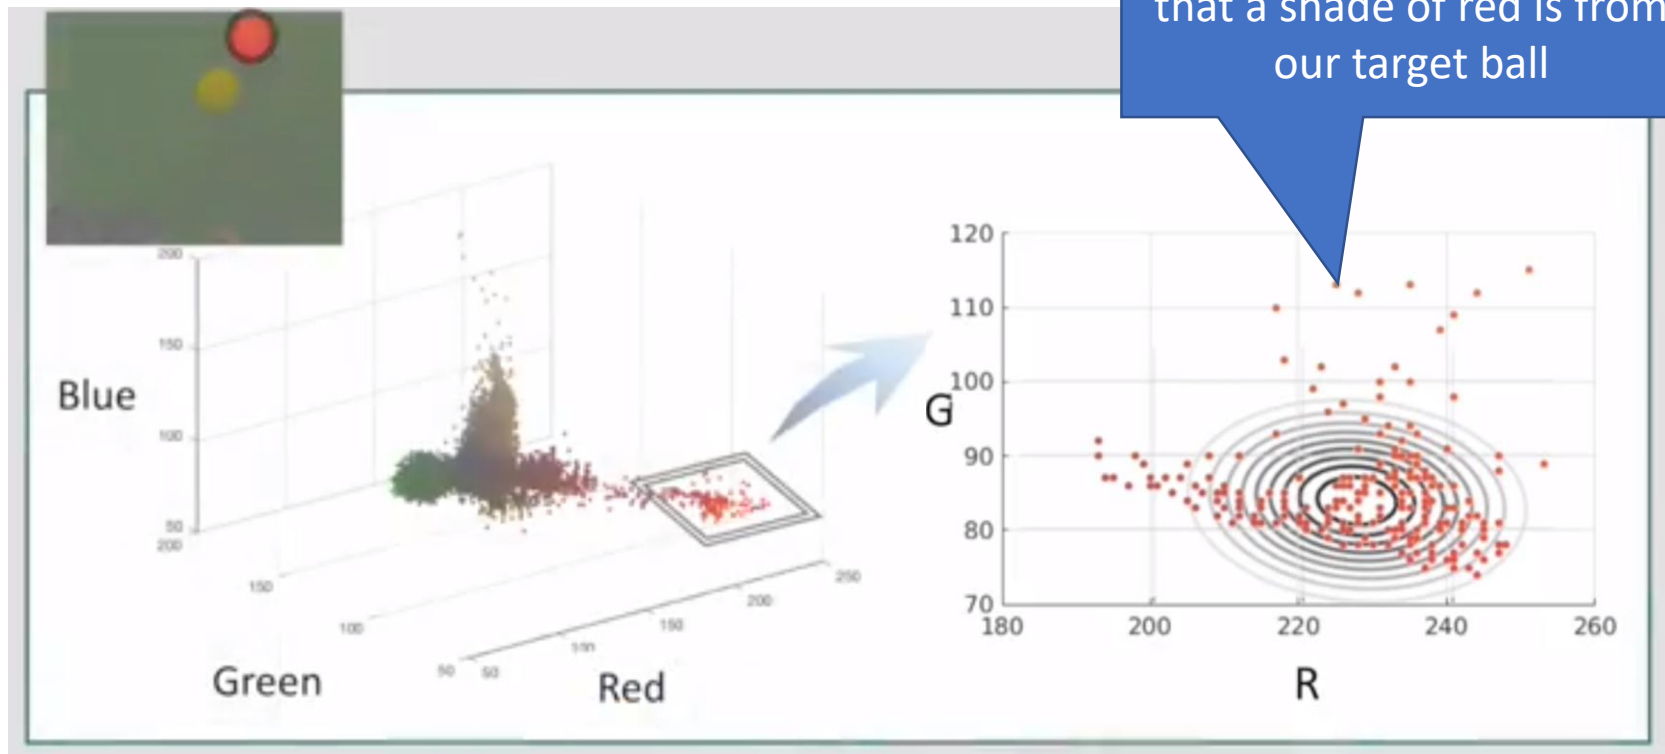
Important: $w_k > 0$, $\sum_{k=1}^K w_k = 1$ must hold

Example: Color Filtering



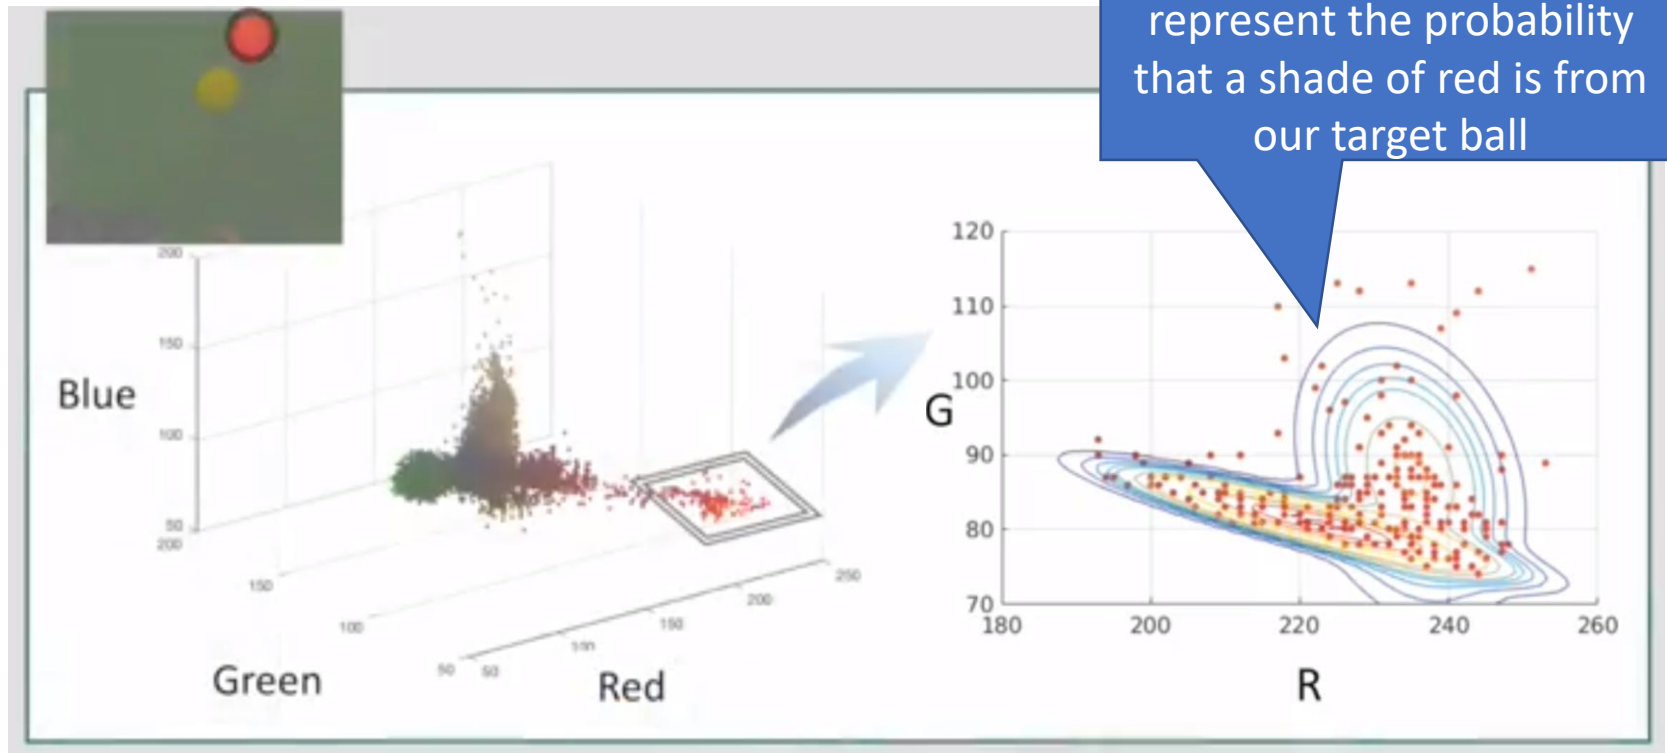
Example: Color Filtering

A single 2D Gaussian can represent the probability that a shade of red is from our target ball



Example: Color Filtering

A mixture of two 2D Gaussians can **better** represent the probability that a shade of red is from our target ball



Exercise:

Improving the original Roomba vacuum cleaner

- Current Operation: Random walk with Vacuum on
 - Limitations:
 - Vacuum operation consumes battery
 - Doesn't return home to charge
 - Has IR sensors to detect ledges so it doesn't fall down stairs
 - Has bump sensors to detect collisions with stationary objects
- If the Roomba had a sensor that could detect dirt underneath it...

How could we design a smarter Roomba?

Objective: Minimize the amount of dirt on the floor at any given moment

Ideas:

- Odometry
- Mapping
- Machine Learning



Hidden Markov Models

Variables:

$$S = s_1, s_2, \dots, s_N$$

(States)

$$V = v_1, v_2, \dots, v_k$$

(Observation Vocab.)

$$A = a_{11}, \dots, a_{ij}, \dots, a_{NN}$$

(Transition prob. Matrix)

$$B = P(o_t | s_i) \forall i \in [1, N], t \in [1, T]$$

(Obs. Emission Probs)

$$\pi = \pi_1, \pi_2, \dots, \pi_N$$

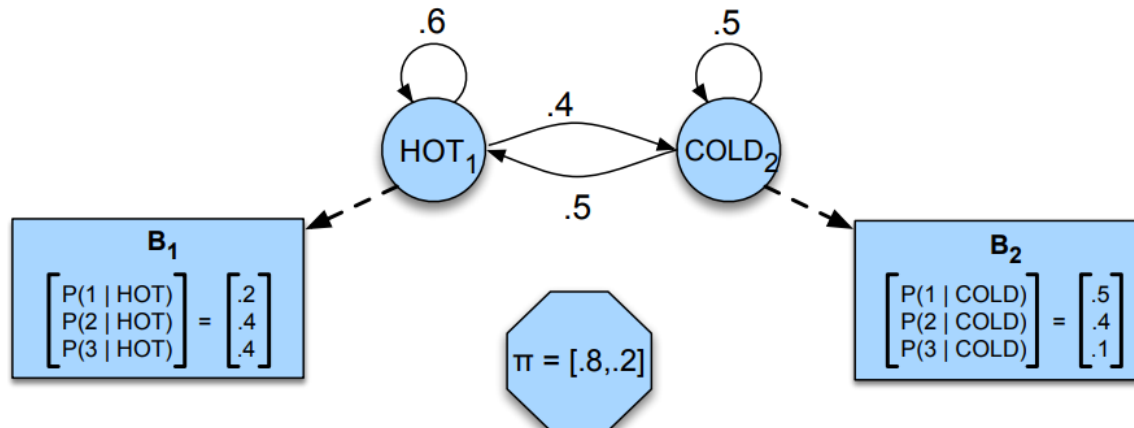
(Initial prob. distribution)

$$O = o_1, o_2, \dots, o_T$$

(Observation Sequence)

$$Q = s_1, s_2, \dots, s_T$$

(State Sequence)



Three Types of Problems

- **Likelihood:**

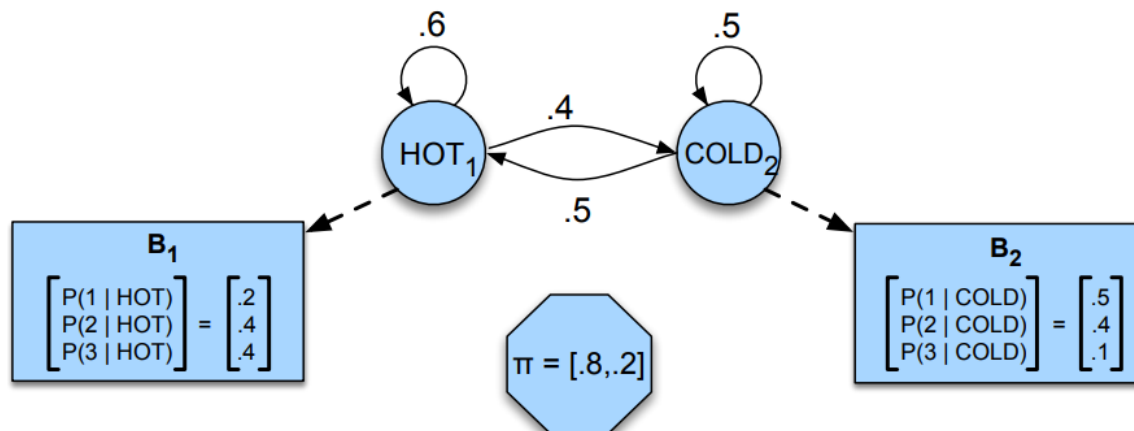
Given $A, B, O \dots$ Determine $P(O|A, B)$

- **Decoding:**

Given $A, B, O \dots$ Determine the 'best' hidden state sequence

- **Learning:**

Given O and $S \dots$ Determine A, B



Likelihood Computation

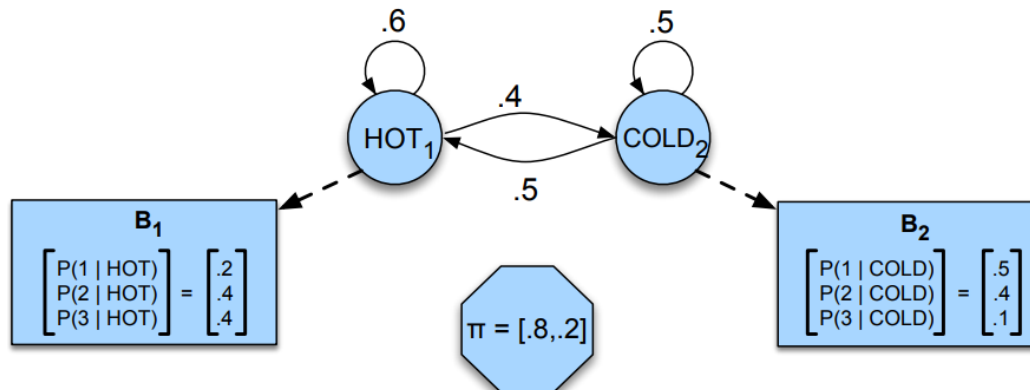
Likelihood: Given A, B, O ... Determine $P(O|A, B)$

Example: $O = \{3, 1, 3\}$

$$P(O|Q) = \prod_{i=1}^T P(o_i|q_i) \text{ -- Prob. of } O \text{ given State Seq. } Q$$

For one possible state sequence (*hot, hot, cold*):

$$P(3 \ 1 \ 3 | \text{hot hot cold}) = P(3|\text{hot}) \times P(1|\text{hot}) \times P(3|\text{cold})$$



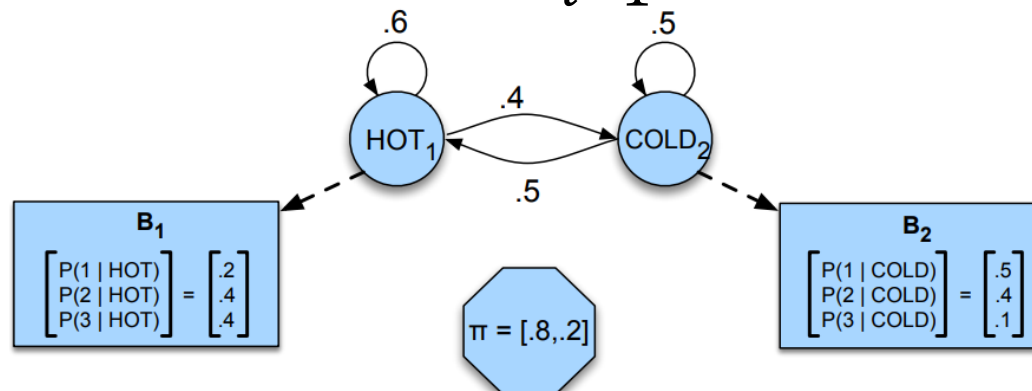
Likelihood Computation

Example: Given A, B and $O = \{3, 1, 3\}$ -- Determine $P(O|A, B)$

But we don't know the state sequence!

Instead, we must weight each sequence by its probability.

$$P(O, Q) = P(O|Q) \times P(Q) = \prod_{i=1}^T P(o_i|q_i) \times \prod_{i=1}^T P(q_i|q_{i-1})$$



Likelihood Computation

$$P(O|Q) = \prod_{i=1}^T P(o_i|q_i)$$

$$P(O, Q) = P(O|Q) \times P(Q) = \prod_{i=1}^T P(o_i|q_i) \times \prod_{i=1}^T P(q_i|q_{i-1})$$

$$P(O) = \sum_Q P(O, Q) = \sum_Q P(O|Q)P(Q)$$



N^T Sequences!

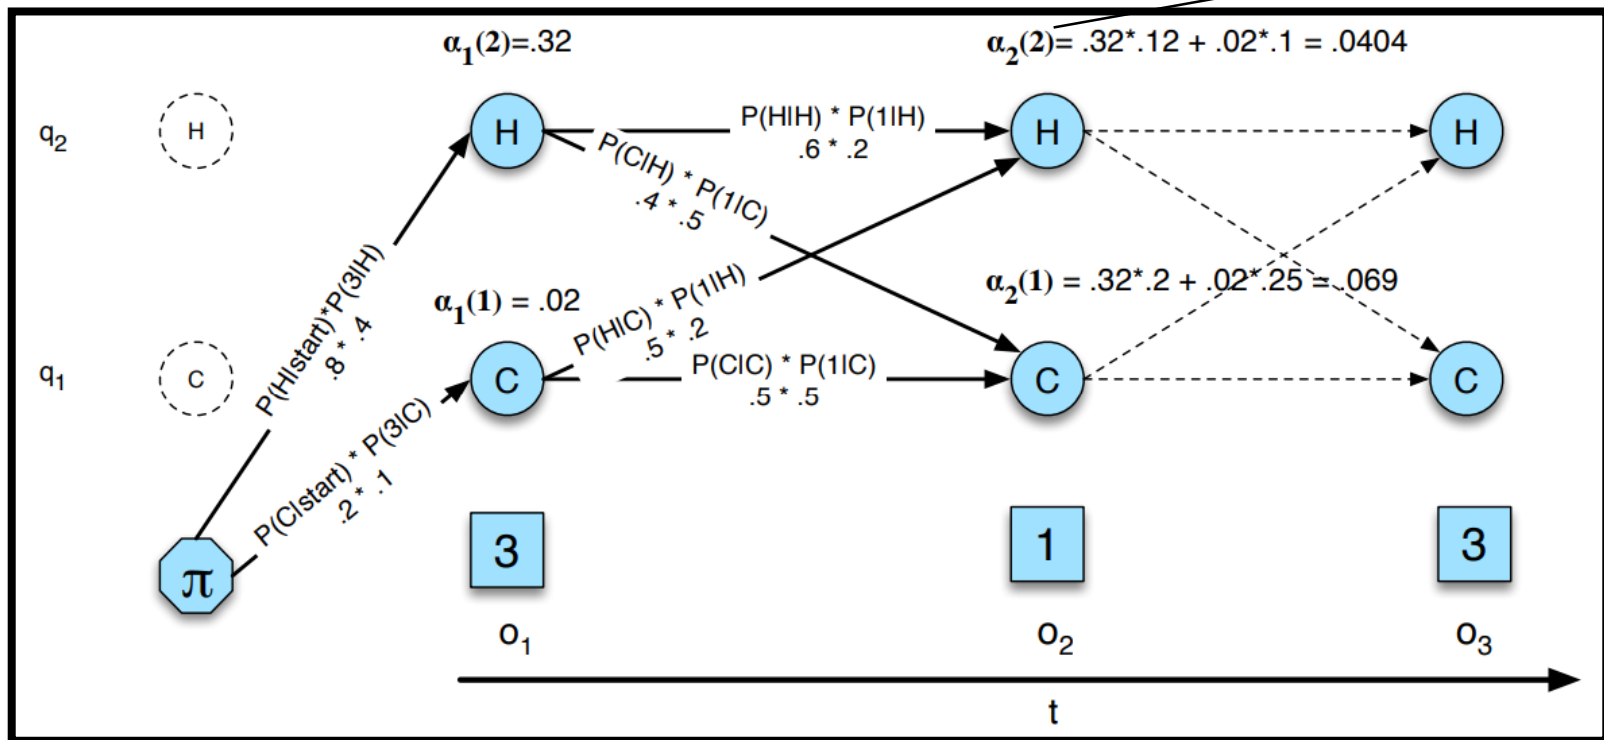
Likelihood Computation: Forward Algorithm

Example: Given A, B and $O = \{3, 1, 3\}$ -- Determine $P(O|A, B)$

Infeasible to solve with $O(N^T)$ algorithm.

Can do it in $O(N^2T)$ with Dynamic Programming!

$$\alpha_t(j): \\ P(s=j|T=t)$$



Likelihood Computation: Forward Algorithm

Example: Given A, B and $O = \{3, 1, 3\}$ -- Determine $P(O|A, B)$

$$\alpha_t(j) = P(o_1, o_2, \dots, o_t, q_t = j \mid A, B)$$

$$\alpha_t(j) = \sum_{i=1}^N \alpha_{t-1}(i) * a_{ij} * b_j(o_t)$$

Prev P(i -> j) P(o|s)

function FORWARD(*observations* of len T , *state-graph* of len N) **returns** *forward-prob*

create a probability matrix *forward*[N, T]

for each state s **from** 1 **to** N **do** ; initialization step

forward[$s, 1$] $\leftarrow \pi_s * b_s(o_1)$

for each time step t **from** 2 **to** T **do** ; recursion step

for each state s **from** 1 **to** N **do**

$$\text{forward}[s, t] \leftarrow \sum_{s'=1}^N \text{forward}[s', t-1] * a_{s', s} * b_s(o_t)$$

forwardprob $\leftarrow \sum_{s=1}^N \text{forward}[s, T]$; termination step

return *forwardprob*

State Sequence Computation: Viterbi Algorithm

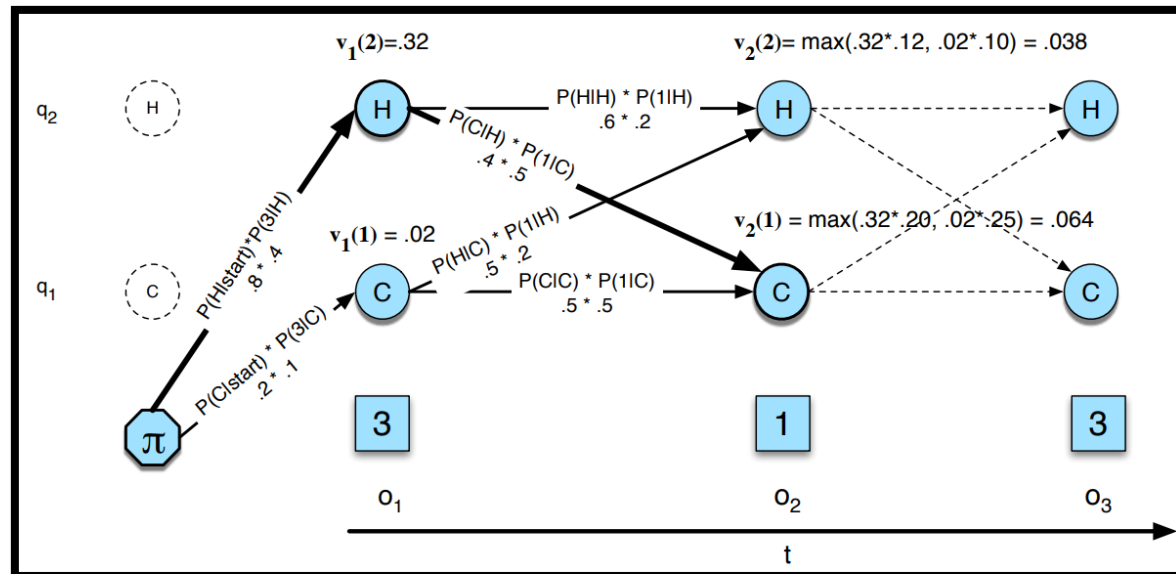
Decoding: Given **A**, **B**, **O** ... determine **Q**

Option 1: Run Forward Algorithm on all state sequences

Option 2: Use Dynamic Programming

$$v_t(j) = \max_{q_1, \dots, q_{t-1}} P(q_1 \dots q_{t-1}, o_1 \dots o_t, q_t = j \mid A, B)$$

$$v_t(j) = \max_{i \in [1, N]} v_{t-1}(i) * a_{ij} * b_j(o_t)$$



State Sequence Computation: Viterbi Algorithm

function VITERBI(*observations* of len T , *state-graph* of len N) **returns** *best-path*, *path-prob*

create a path probability matrix $viterbi[N, T]$

for each state s **from** 1 **to** N **do**

; initialization step

$$viterbi[s, 1] \leftarrow \pi_s * b_s(o_1)$$

$$backpointer[s, 1] \leftarrow 0$$

for each time step t **from** 2 **to** T **do**

; recursion step

for each state s **from** 1 **to** N **do**

$$viterbi[s, t] \leftarrow \max_{s'=1}^N viterbi[s', t-1] * a_{s', s} * b_s(o_t)$$

$$backpointer[s, t] \leftarrow \operatorname{argmax}_{s'=1}^N viterbi[s', t-1] * a_{s', s} * b_s(o_t)$$

$$bestpathprob \leftarrow \max_{s=1}^N viterbi[s, T]$$

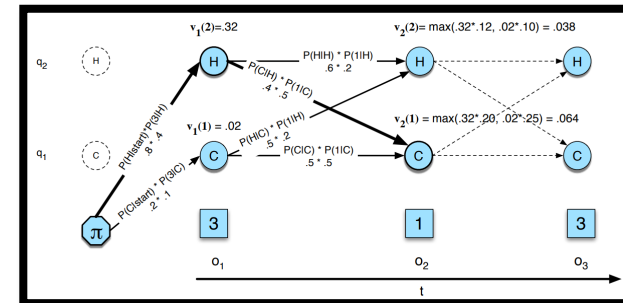
; termination step

$$bestpathpointer \leftarrow \operatorname{argmax}_{s=1}^N viterbi[s, T]$$

; termination step

$bestpath \leftarrow$ the path starting at state $bestpathpointer$, that follows $backpointer[]$ to states back in time

return $bestpath$, $bestpathprob$



Learning an HMM's Parameters

Learning: Given \mathbf{O} and \mathbf{S} ... Determine \mathbf{A}, \mathbf{B}

Challenge: Must simultaneously determine transition probabilities AND emission probabilities!

Special case of Expectation-Maximization, iteratively improving an initial estimate.

But first, let's solve for a Markov Chain (*fully observable*) given $\mathbf{O}, \mathbf{S}, \mathbf{Q}$

Learning Not-so-HMM Parameters A, B

Given Sequences: $\{ 3H, 3H, 2C \} \{ 1C, 1C, 2C \} \{ 1C, 2H, 3H \}$

Compute Initial Probabilities:

$$\pi = \begin{cases} H: 1/3 \\ C: 2/3 \end{cases}$$

Compute Transition Probabilities:

$P(H H) = 2/3$	$P(H C) = 1/2$
$P(C H) = 1/3$	$P(C C) = 1/2$

$$A = \begin{cases} HH: \frac{2}{3} & CH: \frac{1}{3} \\ HC: \frac{1}{2} & CC: \frac{1}{2} \end{cases}$$

Compute Emission Probabilities:

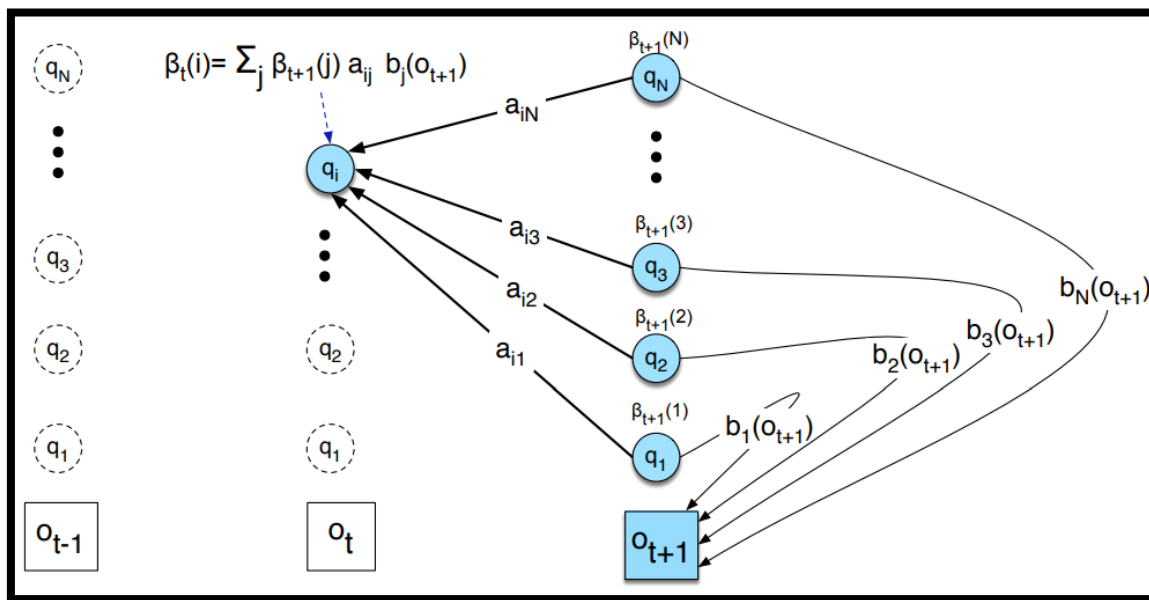
$P(1 \text{hot}) = 0$	$P(1 \text{cold}) = 3/5$
$P(2 \text{hot}) = 1/4$	$P(2 \text{cold}) = 2/5$
$P(3 \text{hot}) = 3/4$	$P(3 \text{cold}) = 0$

$$B = \begin{cases} \text{Hot} = \{ 0, \frac{1}{4}, \frac{3}{4} \} \\ \text{Cold} = \{ \frac{3}{5}, \frac{2}{5}, 0 \} \end{cases}$$

Backward Algorithm

Backward Probability: $\beta_t(i) = P(o_{t+1}, o_{t+2}, \dots o_T | q_t = i, A, B)$

If we're in state i at time t , what's $P(\text{Obs})$ from then to end?



Initialization: $\beta_T(i) = 1, i \in [1, N]$

Recursion: $\beta_t(i) = \sum_{j=1}^N a_{ij} * b_j(o_{t+1}) * \beta_{t+1}(j), i \in [1, N], t \in [1, T)$

Termination: $P(O|A, B) = \sum_{j=1}^N \pi_j * b_j(o_1) * \beta_1(j)$

Forward-Backward: Learning A

$$\hat{a}_{ij} = \frac{\textit{Expected \#transitions from } i \textit{ to } j}{\textit{Expected \#transitions from } i}$$

To compute numerator:

1. Assume we have probability estimate for $i \rightarrow j$ at time t
2. Now assume we had that for all t : sum over all $t \in [0, T)$ to get the total count for $i \rightarrow j$

Define ξ_t as probability of transition from i to j at time t :

$$\xi_t(i, j) = P(q_t = i, q_{t+1} = j \mid O, A, B)$$

...But we don't know the relation between O and Q !

Forward-Backward: Learning A

Define ξ_t as probability of transition from i to j at time t :

$$\xi_t(i, j) = P(q_t = i, q_{t+1} = j \mid O, A, B)$$

...But we don't know the relation between O and Q !

So we define $\text{sort-of-}\xi_t(i, j) = P(q_t = i, q_{t+1} = j, O \mid A, B)$

$$\text{sort-of-}\xi_t(i, j) = \alpha_t(i) * a_{ij} * b_j(o_{t+1}) * \beta_{t+1}(j)$$

Forward

$$\alpha_t(j) = P(o_1, o_2, \dots, o_t, q_t = j \mid A, B)$$

Backward

$$\beta_t(i) = P(o_{t+1}, o_{t+2}, \dots, o_T \mid q_t = i, A, B)$$

Forward-Backward: Learning A

$$\text{sort-of-}\xi_t(i, j) = \alpha_t(i) * a_{ij} * b_j(o_{t+1}) * \beta_{t+1}(j)$$

How do we go from $P(q_t = i, q_{t+1} = j, O | A, B)$ to $P(q_t = i, q_{t+1} = j | O, A, B)$

$$\text{Recall: } P(X|Y, Z) = \frac{P(X, Y|Z)}{P(Y|Z)}$$

Thus, because $P(O|A, B) = \sum_{j=1}^N \alpha_t(j) * \beta_t(j)$

$$\xi_t(i, j) = \frac{\alpha_t(i) * a_{ij} * b_j(o_{t+1}) * \beta_{t+1}(j)}{\sum_{j=1}^N \alpha_t(j) * \beta_t(j)}$$

Forward

$$\alpha_t(j) = P(o_1, o_2, \dots, o_t, q_t = j | A, B)$$


Backward

$$\beta_t(i) = P(o_{t+1}, o_{t+2}, \dots, o_T | q_t = i, A, B)$$

Forward-Backward: Learning A

To compute numerator:

1. Assume we have probability estimate for $i \rightarrow j$ at time t
2. Now assume we had that for all t : sum over all $t \in [0, T)$ to get the total count for $i \rightarrow j$

$$\hat{a}_{ij} = \frac{\text{Expected \#transitions from } i \text{ to } j}{\text{Expected \#transitions from } i} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \sum_{k=1}^N \xi_t(i, k)}$$


$$\xi_t(i, j) = \frac{\alpha_t(i) * a_{ij} * b_j(o_{t+1}) * \beta_{t+1}(j)}{\sum_{j=1}^N \alpha_t(j) * \beta_t(j)}$$

$$\xi_t(i, j) = P(q_t = i, q_{t+1} = j \mid O, A, B)$$

Forward

$$\alpha_t(j) = P(o_1, o_2, \dots, o_t, q_t = j \mid A, B)$$

Backward

$$\beta_t(i) = P(o_{t+1}, o_{t+2}, \dots, o_T \mid q_t = i, A, B)$$

Forward-Backward: Learning B

$$\hat{a}_{ij} = \frac{\text{Expected \#transitions from } i \text{ to } j}{\text{Expected \#transitions from } i} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \sum_{k=1}^N \xi_t(i, k)}$$

$$\xi_t(i, j) = P(q_t = i, q_{t+1} = j \mid O, A, B)$$

Now we need to compute observation emission probability:

$$\hat{b}_j(v_k) = \frac{\text{Expected \# of } v_k \text{ seen in state } j}{\text{Expected \#times in state } j}$$

Forward

$$\alpha_t(j) = P(o_1, o_2, \dots, o_t, q_t = j \mid A, B)$$

Backward

$$\beta_t(i) = P(o_{t+1}, o_{t+2}, \dots, o_T \mid q_t = i, A, B)$$

Forward-Backward: Learning B

Now we need to compute observation emission probability:

$$\hat{b}_j(v_k) = \frac{\text{Expected \# of } v_k \text{ seen in state } j}{\text{Expected \#times in state } j}$$

But first, we need to know **prob. of being in state j at time t**

$$\gamma_t(j) = P(q_t = j \mid O, A, B) = \frac{P(q_t = j, O \mid A, B)}{P(O \mid A, B)}$$

$$\gamma_t(j) = \frac{\alpha_t(j) * \beta_t(j)}{P(O \mid A, B)}$$

Forward

$$\alpha_t(j) = P(o_1, o_2, \dots, o_t, q_t = j \mid A, B)$$

Backward

$$\beta_t(i) = P(o_{t+1}, o_{t+2}, \dots, o_T \mid q_t = i, A, B)$$

Forward-Backward: Learning B

Now we need to compute observation emission probability:

$$\hat{b}_j(v_k) = \frac{\text{Expected \# of } v_k \text{ seen in state } j}{\text{Expected \#times in state } j} = \frac{\sum_{t=1}^T \gamma_t(j) * I(o_t = v_k)}{\sum_{t=1}^T \gamma_t(j)}$$

$\gamma_t(j)$ = prob. of being in state j at time t

$$\gamma_t(j) = \frac{\alpha_t(j) * \beta_t(j)}{P(O|A, B)}$$

Forward

$$\alpha_t(j) = P(o_1, o_2, \dots, o_t, q_t = j \mid A, B)$$

Backward

$$\beta_t(i) = P(o_{t+1}, o_{t+2}, \dots, o_T \mid q_t = i, A, B)$$

Expectation-Maximization on A, B

E-Step: Compute state occupancy count γ , expected state transition count ξ using existing A, B probabilities

M-Step: Compute A, B using existing γ and ξ probabilities

$\alpha_t(j)$ = prob. to be in state j at t

$\beta_t(j)$ = prob. of O from state j at t

$\xi_t(i, j)$ = prob. of transition from i to j at time t

$\gamma_t(j)$ = prob. of being in state j at time t

function FORWARD-BACKWARD(*observations* of len T , *output vocabulary* V , *hidden state set* Q) **returns** $HMM=(A, B)$

initialize A and B

iterate until convergence

E-step

$$\gamma_t(j) = \frac{\alpha_t(j)\beta_t(j)}{\alpha_T(q_F)} \quad \forall t \text{ and } j$$

$$\xi_t(i, j) = \frac{\alpha_t(i)a_{ij}b_j(o_{t+1})\beta_{t+1}(j)}{\alpha_T(q_F)} \quad \forall t, i, \text{ and } j$$

M-step

$$\hat{a}_{ij} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \sum_{k=1}^N \xi_t(i, k)}$$

$$\hat{b}_j(v_k) = \frac{\sum_{t=1 \text{ s.t. } O_t=v_k}^T \gamma_t(j)}{\sum_{t=1}^T \gamma_t(j)}$$

return A, B