The background is a collage of images related to human-robot interaction. It includes a person in a black shirt and yellow gloves interacting with a robot arm, a person in a white shirt interacting with a robot arm, a person in a black shirt interacting with a robot arm, a person in a white shirt interacting with a robot arm, a person in a black shirt interacting with a robot arm, and a person in a white shirt interacting with a robot arm.

Algorithmic Human-Robot Interaction

Reinforcement Learning

CSCI 7000

Prof. Brad Hayes

Computer Science Department

University of Colorado Boulder

Final Papers

Due Tuesday at 11:59pm

- No presentation required – Online submission only

Mandatory Files

- Zip file titled “Paper.zip”
 - Final PDF
 - LaTeX / Word files
 - Figures
 - Link to a video presentation (up to 10min)
 - Narration over slides is preferred
- Zip file titled “Code.zip”
 - Checkout of your (up-to-date) Git repository
 - README file including instructions to run your project

Highly Encouraged

- Video demonstration of results in your presentation

Paper for Today:

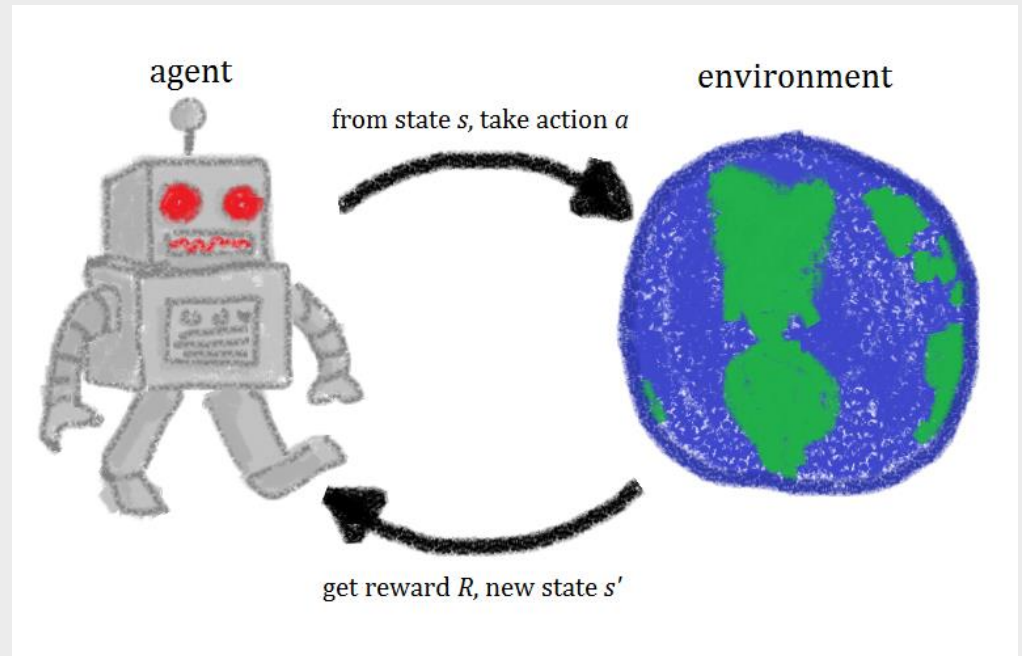
A Reduction of Imitation Learning and Structured Prediction to No-Regret Online Learning

by Ross et al.

Pro: Nishank Sharma

Con: Ashwin Vasan

Getting Started with RL



Prerequisites

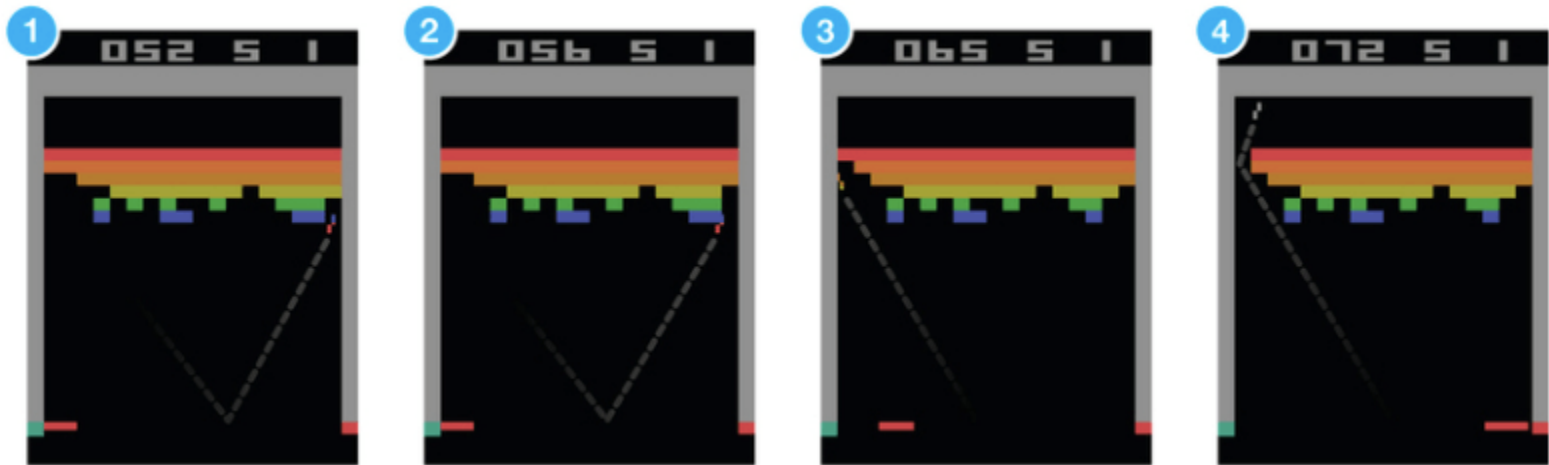
Numpy: `pip install numpy`

OpenAI Gym: `pip install gym`

Torch: `pip install torch`

Download the `rl-demo.py` source code on Moodle!

Reinforcement Learning



Unsupervised Learning

- No labels at all!
- Only given data
(Clustering/embedding problem)

Reinforcement Learning

Sparse labeling
(Few states get a correct answer)

Time-delayed signaling
(Labels not provided in a timely manner)

Supervised Learning

Target label for every example
(each state has the 'correct' action)

No need for reward function!

Breakout:

State: Position of ball, paddle, bricks, etc.

Observation: Pixels from screen

Actions: Left, Right, Release Ball

Exploration / Exploitation

How often should you listen to your own strategy?

When should you give up acting randomly to find a new strategy?

A typical formulation is ϵ -greedy exploration:

Follow policy π most of the time ($\epsilon\%$), act randomly $(1 - \epsilon)\%$ of the time

More complex functions can also work well:

- Set an ϵ schedule (start value α , end value ω , decay δ)
- $\epsilon = \omega + \max(0, (\alpha - \omega)) * e^{-t/\delta}$

Discounted Future Reward

To perform well long-term, the agent needs to consider immediate rewards AND future rewards:

$$R = r_1 + r_2 + \cdots + r_n$$

But stochastic environments don't make r_{t+1} a certainty

The further into the future we look, the less certain we can be.

We address this by **discounting future reward**:

$$R = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \cdots + \gamma^{n-t} r_n$$

Discounted Future Reward \rightarrow Q-Function

$$R_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \cdots + \gamma^{n-t} r_n$$

$$R_t = r_t + \gamma(r_{t+1} + \gamma(r_{t+2} + \cdots))$$

$$R_t = r_t + \gamma R_{t+1}$$

$$Q(s_t, a_t) = \max R_{t+1}$$

Q tells us the best possible score at game-end if I do a_t in state s_t

$$\pi(s) = \operatorname{argmax}_a Q(s, a)$$

Deriving the Q-Function

$$Q(s_t, a_t) = \max R_{t+1}$$

Q tells us the best possible score at game-end if I do a_t in state s_t

$$\pi(s) = \operatorname{argmax}_a Q(s, a)$$

How do we derive Q?

Consider a single transition:

$\langle s, a, r, s' \rangle ::$ (state, action, reward, next state)

$$Q(s, a) = r + \gamma * \max_{a'} Q(s', a')$$

```
initialize  $Q[num\_states, num\_actions]$  arbitrarily
observe initial state  $s$ 
repeat
    select and carry out an action  $a$ 
    observe reward  $r$  and new state  $s'$ 
     $Q[s, a] = Q[s, a] + \alpha(r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$ 
     $s = s'$ 
until terminated
```

Q-Learning: Basic Algorithm

$$Q(s, a) = r + \gamma * \max_{a'} Q(s', a')$$

```

initialize  $Q[num\_states, num\_actions]$  arbitrarily
observe initial state  $s$ 
repeat
    select and carry out an action  $a$ 
    observe reward  $r$  and new state  $s'$ 
     $Q[s, a] = Q[s, a] + \alpha(r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$ 
     $s = s'$ 
until terminated

```

```

self.Q = np.random.random([NUM_BINS, NUM_BINS, NUM_BINS, NUM_BINS, 2]) # Value of st

def select_action(self, state):
    sample = random.random()
    eps_threshold = self.EPS_END + (self.EPS_START - self.EPS_END)
        * math.exp(-1. * self.steps_done / self.EPS_DECAY)
    self.steps_done += 1
    if sample > eps_threshold: # Exploit vs. Explore check
        return self.Q[tuple(state)].argmax()
    else:
        return random.randrange(2) # Pick random action

def update_model(self, state, action, reward, state_next):
    state_action_q_index = tuple(np.hstack([state, action]))
    max_future_q = self.Q[tuple(state_next)].max()
    self.Q[state_action_q_index] = self.Q[state_action_q_index] + self.LEARNING_RATE
        * (reward + self.GAMMA * max_future_q - self.Q[state_action_q_index])

```

Tabular Q-Learning Doesn't Scale

DeepMind used 4 grayscale 84x84 frames for their Nature paper
“Human-level control through deep reinforcement learning”

State Space Issues:

- 84x84 Pixel Screen * 4 frames @ 256 grayscale levels = $256^{84*84*4} \approx 10^{67970}$
- Q-Table would have to have 10^{67970} rows!
- Most states are never visited, others very rarely

Solution:

- Need a function approximator that's very good at learning features for structured data
- Deep Learning to the rescue

Deep Q-Network: Neural Net as Q-Function

Deep Q-Network Architecture

Layer	Input	Filter size	Stride	Num filters	Activation	Output
conv1	84x84x4	8x8	4	32	ReLU	20x20x32
conv2	20x20x32	4x4	2	64	ReLU	9x9x64
conv3	9x9x64	3x3	1	64	ReLU	7x7x64
fc4	7x7x64			512	ReLU	512
fc5	512			18	Linear	18

(Convolutional Neural Network architecture from Mnih et al.)

18 possible output actions

Deep Q-Network Architecture + Loss

Layer	Input	Filter size	Stride	Num filters	Activation	Output
conv1	84x84x4	8x8	4	32	ReLU	20x20x32
conv2	20x20x32	4x4	2	64	ReLU	9x9x64
conv3	9x9x64	3x3	1	64	ReLU	7x7x64
fc4	7x7x64			512	ReLU	512
fc5	512			18	Linear	18

$$L = \frac{1}{2} \left[\underbrace{r + \max_{a'} Q(s', a')}_{\text{Target}} - \underbrace{Q(s, a)}_{\text{Prediction}} \right]^2$$

Experience Replay

- Deep networks approximate highly non-linear functions
- Updates at each time step are too similar
 - Drives network into a local minimum
 - Need to diversify updates to prevent this from happening
- Solution:
 - Store all $\langle s, a, r, s' \rangle$ tuples experienced during learning
 - Pick random batches of experiences to train with at each update step
 - Makes sure your network doesn't “forget” previous information



Deep Q-Learning

```
initialize replay memory  $D$ 
initialize action-value function  $Q$  with random weights
observe initial state  $s$ 
repeat
    select an action  $a$ 
        with probability  $\epsilon$  select a random action
        otherwise select  $a = \operatorname{argmax}_{a'} Q(s, a')$ 
    carry out action  $a$ 
    observe reward  $r$  and new state  $s'$ 
    store experience  $\langle s, a, r, s' \rangle$  in replay memory  $D$ 

    sample random transitions  $\langle ss, aa, rr, ss' \rangle$  from replay memory  $D$ 
    calculate target for each minibatch transition
        if  $ss'$  is terminal state then  $tt = rr$ 
        otherwise  $tt = rr + \gamma \max_{a'} Q(ss', aa')$ 
    train the  $Q$  network using  $(tt - Q(ss, aa))^2$  as loss

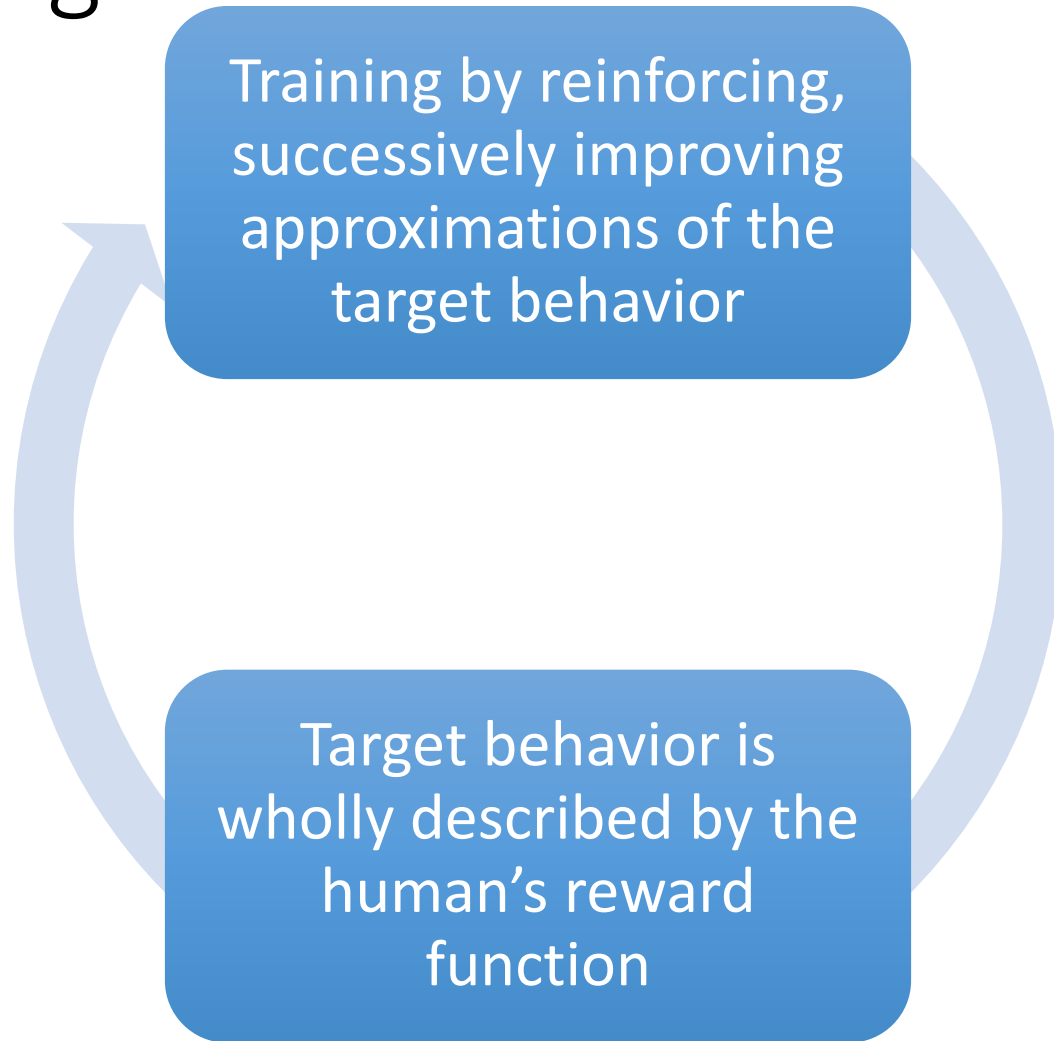
     $s = s'$ 
until terminated
```



Search is bad at forecasting



Shaping



Why Bother “Shaping”?


- Decreases sample complexity
- Reduces or eliminates dependence on environment reward
- Doesn't require a domain expert



Shaping: Iteratively improving agent policy using a human reward signal

Related Work

- Learning from Advice
 - Works for people, but extremely complicated
- Learning from Demonstration
 - Human performance replaces standard environmental reward
 - Limited:
 - Human must be an expert
 - Not always a clear interface for the human to demonstrate with



TAMER Framework

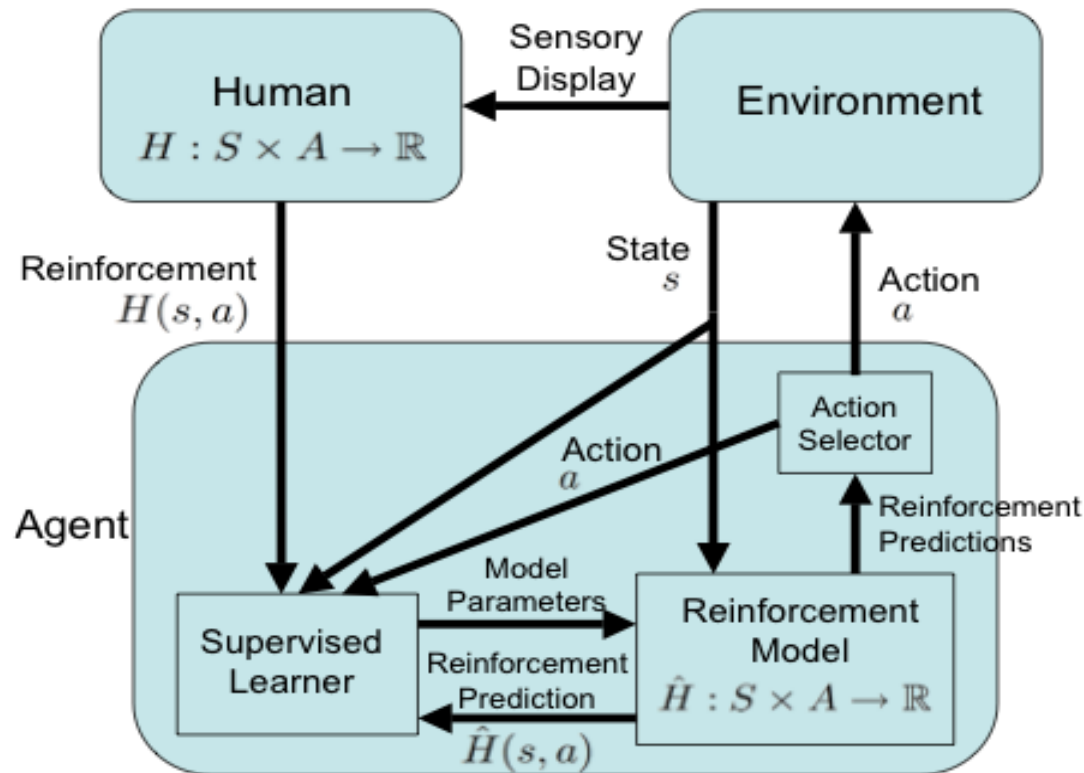


Figure 1: Framework for Training an Agent Manually via Evaluative Reinforcement (TAMER).

Human
automatically
engages foresight

Algorithm doesn't
have to look ahead
quite as far

Human can act in
a tight temporal
window

Credit assignment can
be frequent and be
indicative of long term
benefit to
state-action pairs

Advantages of Human Signal

Human reward is a moving target

- Human reinforcement will be more sparse over time
- Complacency causes agent to level performance

Challenges of Human Signal

Exploration

- Open research problem
- “To be filled by the agent designer”
- Experiments suggest that greedy selection provides sufficient exploration
 - Traditional RL maximizes return
 - Discounted sum of all future reward
 - TAMER maximizes immediate return without regard for future states

Exploration Continued

- TAMER acts in the present
 - Assumes the human is thinking ahead for the agent
 - Becomes problematic if sensors exceed human perception of the environment
 - Vehicle is moving too quickly for human to perceive nail in the road
 - Glint detector would detect object
 - UAV radar detects an inbound missile
 - Human can't react quickly enough to change agent behavior



Credit Assignment

- Time step frequency can exceed human capacity for response
 - Anything below 200ms is too fast
- Credit provided via simple algorithm:
 - Choose PDF (Gamma distribution)
 - Probability over time of particular time slice being targeted with reward
 - Proportion of credit equal to the integral over the execution window between timesteps.
 - Credit used to weight the correction term in Q-function updates

$$\Gamma(\theta, k) = x^{k-1} \frac{\exp(-x/\theta)}{\Gamma(k) \theta^k}$$

Credit Assignment




Figure 3: Probability density function $f(x)$ for a $\text{gamma}(2.0, 0.28)$ distribution. Reinforcement signal h is received at time 0. If t and t' are times of consecutive time steps, credit for the time step at t is $\int_{t'}^t f(x)dx$. Note that time moves backwards as one moves right along the x-axis.

Example

Algorithm 2 A greedy TAMER algorithm with credit assignment, using a linear model and gradient descent updates


Require: *Input: stepSize, windowSize,*

- 1: *Crediter.init(windowSize)*
- 2: $\vec{s} \leftarrow \vec{0}$
- 3: $\vec{f} \leftarrow \vec{0}$
- 4: $\vec{w} \leftarrow \vec{0}$
- 5: **while** true **do**
- 6: *Crediter.updateTime(clockTime())*
- 7: $h \leftarrow \text{getHumanReinfSincePreviousTimeStep}()$
- 8: **if** $h \neq 0$ **then**
- 9: $\vec{\text{credFeats}} \leftarrow \vec{0}$
- 10: **for all** $(f_t, t) \in \text{Crediter.historyWindow}$ **do**
- 11: $c_t \leftarrow \text{Crediter.assignCredit}(t)$
- 12: $\vec{\text{credFeats}} \leftarrow \vec{\text{credFeats}} + (c_t \times \vec{f}_t)$
- 13: **end for**
- 14: $\text{error} \leftarrow h - (\vec{w} \cdot \vec{\text{credFeats}})$
- 15: $\vec{w} \leftarrow \vec{w} + (\text{stepSize} \times \text{error} \times \vec{\text{credFeats}})$
- 16: **end if**
- 17: $\vec{s} \leftarrow \text{getStateVec}()$
- 18: $a \leftarrow \text{argmax}_a (\vec{w} \cdot (\text{getFeatures}(\vec{s}, a)))$
- 19: $\vec{f} \leftarrow \text{getFeatures}(\vec{s}, a)$
- 20: *takeAction(a)*
- 21: *Crediter.updateWindow(\vec{f})*
- 22: wait for next time step
- 23: **end while**

Example Domains: Tetris

- Tetris


- Low update frequency
- Complex problem domain
 - Massive state space
 - TD-Learning struggles
 - Successful implementations cheat
- Stochastic environment contributes to poor performance on other learners
- TAMER learns ~65 lines per game in 3 training episodes
 - Policy search algorithms do much better, but require many , many more training episodes (~1000s of lines cleared)




TAMER in Action



Pre-Training



Training



Post-Training

Tetris

Table 1: Results of various Tetris agents.

Method	Mean Lines Cleared		Games for Peak
	at Game 3	at Peak	
TAMER	65.89	65.89	3
RRL-KBR [15]	5	50	120
Policy Iteration [2]	~ 0 (no learning until game 100)	3183	1500
Genetic Algorithm [5]	~ 0 (no learning until game 500)	586,103	3000
CE+RL [17]	~ 0 (no learning until game 100)	348,895	5000